

NOVA Microhypervisor Interface Specification

Udo Steinberg
udo@hypervisor.org

July 21, 2021

Preliminary

Preliminary

Copyright © 2006–2011 Udo Steinberg, Technische Universität Dresden

Copyright © 2012–2013 Udo Steinberg, Intel Corporation

Copyright © 2014–2016 Udo Steinberg, FireEye, Inc.

Copyright © 2019–2021 Udo Steinberg, BedRock Systems, Inc.

This specification is provided "as is" and may contain defects or deficiencies which cannot or will not be corrected. The author makes no representations or warranties, either expressed or implied, including but not limited to, warranties of merchantability, fitness for a particular purpose, or non-infringement that the contents of the specification are suitable for any purpose or that any practice or implementation of such contents will not infringe any third party patents, copyrights, trade secrets or other rights.

The specification could include technical inaccuracies or typographical errors. Additions and changes are periodically made to the information therein; these will be incorporated into new versions of the specification, if any.

Contents

I	Introduction	1
1	System Architecture	2
II	Basic Abstractions	3
2	Kernel Objects	4
2.1	Protection Domain	4
2.1.1	Object Space	4
2.1.2	Memory Space	4
2.1.3	I/O Port Space	4
2.1.4	MSR Space	4
2.2	Execution Context	5
2.3	Scheduling Context	5
2.4	Portal	5
2.5	Semaphore	6
III	Application Programming Interface	7
3	Data Types	8
3.1	Capability	8
3.1.1	Null Capability	8
3.1.2	Object Capability	8
3.1.2.1	PD Object Capability	8
3.1.2.2	EC Object Capability	8
3.1.2.3	SC Object Capability	8
3.1.2.4	PT Object Capability	9
3.1.2.5	SM Object Capability	9
3.1.3	Memory Capability	9
3.1.4	I/O Port Capability	9
3.1.5	MSR Capability	9
3.2	Capability Selector	10
3.3	User Thread Control Block	11
3.3.1	Regular Layout	11
3.3.2	Architectural Layout	11
3.4	Message Transfer Descriptor	12
3.4.1	Regular IPC	12
3.4.2	Architectural IPC	12
4	Hypercalls	13
4.1	Definitions	13
4.1.1	Hypercall Numbers	13
4.1.2	Status Codes	13
4.1.3	Space Type	14
4.1.4	Access Type	14
4.2	Communication	15
4.2.1	IPC Call	15
4.2.2	IPC Reply	16

4.3	Object Creation	17
4.3.1	Create Protection Domain	17
4.3.2	Create Execution Context	18
4.3.3	Create Scheduling Context	20
4.3.4	Create Portal	21
4.3.5	Create Semaphore	22
4.4	Object Control	23
4.4.1	Control Protection Domain	23
4.4.2	Control Execution Context	25
4.4.3	Control Scheduling Context	26
4.4.4	Control Portal	27
4.4.5	Control Semaphore	28
4.5	Platform Management	29
4.5.1	Control Power Management	29
4.5.2	Assign Interrupt	30
4.5.3	Assign Device	32
5	Booting	33
5.1	Microhypervisor	33
5.1.1	ELF Image Loading	33
5.1.2	Platform Resource Access	33
5.2	Root Protection Domain	34
5.2.1	ELF Image Format	34
5.2.2	Initial Configuration	34
5.2.2.1	Object Space	34
5.2.2.2	Memory Space	35
5.3	Hypervisor Information Page	36
IV	Application Binary Interface	38
6	ABI aarch64	39
6.1	Boot State	39
6.1.1	NOVA Microhypervisor	39
6.1.1.1	Multiboot v2 Launch	39
6.1.1.2	Multiboot v1 Launch	39
6.1.1.3	Legacy Launch	39
6.1.2	Root Protection Domain	40
6.2	Protected Resources	41
6.2.1	Memory Space	41
6.3	Physical Memory	41
6.3.1	Memory Map	41
6.4	Virtual Memory	41
6.4.1	Cacheability Attributes	41
6.4.2	Shareability Attributes	41
6.5	Event-Specific Capability Selectors	42
6.5.1	Architectural Events	42
6.5.2	Microhypervisor Events	43
6.6	Architecture-Dependent Structures	44
6.6.1	Hypervisor Information Page	44
6.6.2	User Thread Control Block	45
6.6.3	Message Transfer Descriptor	46
6.7	Calling Convention	47
6.8	Supplementary Functionality	49

7	ABI x86-64	51
7.1	Boot State	51
7.1.1	NOVA Microhypervisor	51
7.1.1.1	Multiboot v2 Launch	51
7.1.1.2	Multiboot v1 Launch	51
7.1.2	Root Protection Domain	51
7.2	Protected Resources	52
7.2.1	Memory Space	52
7.2.2	I/O Port Space	52
7.3	Physical Memory	52
7.3.1	Memory Map	52
7.4	Virtual Memory	52
7.4.1	Cacheability Attributes	52
7.4.2	Shareability Attributes	52
7.5	Event-Specific Capability Selectors	53
7.5.1	Architectural Events	53
7.5.2	Microhypervisor Events	54
7.6	Architecture-Dependent Structures	55
7.6.1	Hypervisor Information Page	55
7.6.2	User Thread Control Block	55
7.6.2.1	Encoding: Segment Access Rights	56
7.6.2.2	Encoding: Interruption Information	56
7.6.3	Message Transfer Descriptor	57
7.7	Calling Convention	58
V	Appendix	62
A	Acronyms	63
B	Bibliography	65
C	Console	66
C.1	Memory-Buffer Console	66
C.2	UART Console	66
D	Download	67

Notation

Throughout this document, the following symbols are used:

- ~ Indicates that the value of this parameter or field is **undefined**. Future versions of this specification may define a meaning for the parameter or field.
- Indicates that the value of this parameter or field is **ignored**. Future versions of this specification may define a meaning for the parameter or field.
- ≡ Indicates that the value of this parameter or field is **unchanged**. The microhypervisor will preserve the value across hypercalls.

Preliminary

Part I

Introduction

Preliminary

1 System Architecture

The NOVA OS Virtualization Architecture facilitates the coexistence of multiple legacy guest operating systems and a multi-server user-mode framework on a single platform [9]. The core system leverages virtualization technology provided by modern x86 or ARM platforms and comprises the NOVA microhypervisor and one or more Virtual-Machine Monitors (VMMs).

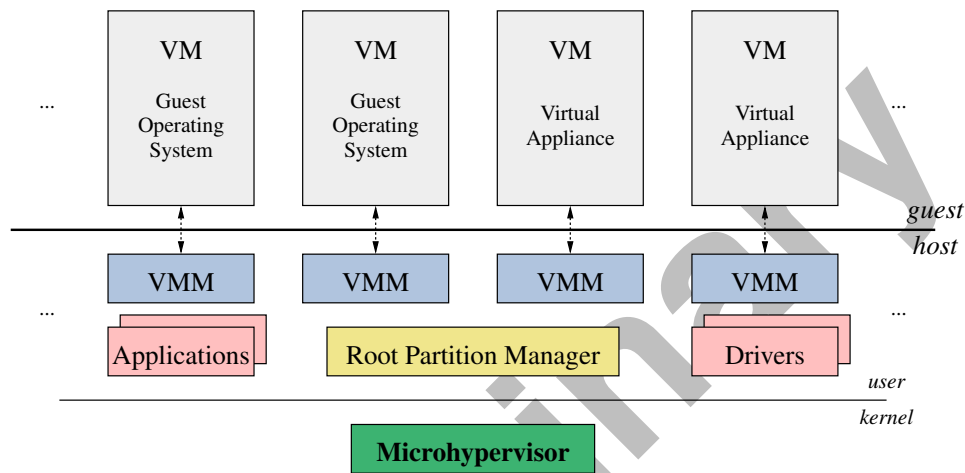


Figure 1.1: System Architecture

Figure 1.1 shows the structure of the system. The microhypervisor is the only component running in privileged root/kernel mode. It isolates the user-level servers, including the virtual-machine monitor, from one another by placing them in different address spaces in unprivileged root/user mode. Each legacy guest operating system runs in its own virtual-machine environment in non-root mode and is therefore isolated from the other components.

Besides isolation, the microhypervisor also provides mechanisms for partitioning and delegation of platform resources, such as CPU time, physical memory, I/O ports and hardware interrupts and for establishing communication paths between different protection domains.

The virtual-machine monitor handles virtualization faults and implements virtual devices that enable legacy guest operating systems to function in the same manner as they would on bare hardware. Providing this functionality outside the microhypervisor in the VMM considerably reduces the size of the trusted computing base for all applications that do not require virtualization support.

The architecture and interfaces of the VMM and the multi-server user-mode framework are not described in this document.

Part II

Basic Abstractions

Preliminary

2 Kernel Objects

2.1 Protection Domain

1. The **Protection Domain (PD)** is a unit of protection and isolation.
2. Access to a **Protection Domain (PD)** is controlled by a **PD Object Capability (CAP_{OBJPD})**.
3. A **PD** is composed of a set of spaces that store **Capabilities (CAPs)** to kernel objects or platform resources that can be accessed by **ECs** within that **PD**. The following subsections detail these spaces.

2.1.1 Object Space

1. The **Object Space (SPC_{OBJ})** is available on all architectures.
2. Each slot of the **Object Space** contains either a **Null Capability (CAP₀)** or an **Object Capability (CAP_{OBJ})** that refers to a kernel object.
3. An **Object Capability Selector (SEL_{OBJ})** serves as index into the **Object Space** and selects a slot.
4. Each hypercall issued from within the **PD** explicitly specifies the **SEL_{OBJ}** to select the **CAP_{OBJ}** for the kernel object on which it operates.

2.1.2 Memory Space

1. The **Memory Space (SPC_{MEM})** is available on all architectures.
2. Each slot of the **Memory Space** contains either a **Null Capability (CAP₀)** or a **Memory Capability (CAP_{MEM})** that refers to a 4 KiB page frame in physical memory.
3. A **Memory Capability Selector (SEL_{MEM})** serves as index into the **Memory Space** and selects a slot.
4. Each memory access issued from within the **PD** implicitly uses the virtual page number (**VirtAddr** >> 12) of the access as **SEL_{MEM}** to select the **CAP_{MEM}** for the 4 KiB page frame on which it operates.

2.1.3 I/O Port Space

1. The **I/O Port Space (SPC_{PtIO})** is available on the x86 architecture only.
2. Each slot of the **I/O Port Space** contains either a **Null Capability (CAP₀)** or an **I/O Port Capability (CAP_{PtIO})** that refers to the physical I/O port corresponding to the slot number.
3. An **I/O Port Capability Selector (SEL_{PtIO})** serves as index into the **I/O Port Space** and selects a slot.
4. Each I/O access (IN/OUT instruction) issued from within the **PD** implicitly uses the I/O port number of the access as **SEL_{PtIO}** to select the **CAP_{PtIO}** for the I/O port on which it operates.

2.1.4 MSR Space

1. The **MSR Space (SPC_{MSR})** is available on the x86 architecture only.
2. Each slot of the **MSR Space** contains either a **Null Capability (CAP₀)** or an **MSR Capability (CAP_{MSR})** that refers to the physical **MSR** corresponding to the slot number.
3. An **MSR Capability Selector (SEL_{MSR})** serves as index into the **MSR Space** and selects a slot.
4. Each **MSR** access (RDMSR/WRMSR instruction) issued from within the **PD** implicitly uses the **MSR** number of the access as **SEL_{MSR}** to select the **CAP_{MSR}** for the **MSR** on which it operates.

2.2 Execution Context

1. The **Execution Context (EC)** is an abstraction for an activity within a **PD**.
2. Access to an **Execution Context (EC)** is controlled by an **EC Object Capability** ($CAP_{OBJ_{EC}}$).
3. An **EC** is permanently bound to exactly one physical CPU.
4. An **EC** is permanently bound to the **PD** for which it was created.
5. There exist three types of **Execution Context**:
 - Local Thread – these may optionally have **PTs** (but not **SCs**) bound to it.
 - Global Thread – these may optionally have an **SC** (but not **PTs**) bound to it.
 - Virtual CPUs – these may optionally have an **SC** (but not **PTs**) bound to it.
6. An **EC** comprises the following state:
 - Reference to bound **PD** (2.1)
 - Event Selector Base (SEL_{EVT})
 - **User Thread Control Block (ARM, x86) (UTCB)** (3.3)
 - Central Processing Unit (**CPU**) registers (architecture dependent)
 - Floating Point Unit (**FPU**) registers (architecture dependent)

2.3 Scheduling Context

1. The **Scheduling Context (SC)** is a unit of dispatching and prioritization.
2. Access to a **Scheduling Context (SC)** is controlled by an **SC Object Capability** ($CAP_{OBJ_{SC}}$).
3. An **SC** is permanently bound to exactly one physical CPU.
4. An **SC** is permanently bound to the **EC** for which it was created.
5. Donation allows another **EC** to consume the budget of the **SC** for the duration of the donation.
6. A scheduling context comprises the following state:
 - Reference to bound **EC** (2.2)
 - Scheduling priority – numerically higher priorities always preempt numerically lower priorities
 - Scheduling budget – time after which the **SC** can be preempted by an **SC** with the same priority

2.4 Portal

1. A **Portal (PT)** represents a dedicated entry point into the **PD** in which the portal was created.
2. Access to a **Portal (PT)** is controlled by a **PT Object Capability** ($CAP_{OBJ_{PT}}$).
3. A **PT** is permanently bound to the **EC** for which it was created.
4. A portal comprises the following state:
 - Reference to bound **EC** (2.2)
 - **Message Transfer Descriptor (ARM, x86) (MTD)** (3.4)
 - Entry instruction pointer
 - Portal Identifier (**PID**)

2.5 Semaphore

1. A [Semaphore \(SM\)](#) provides a means to synchronize execution and interrupt delivery by selectively blocking and unblocking execution contexts.
2. Access to a [Semaphore \(SM\)](#) is controlled by a [SM Object Capability \(CAP_{OBJ_{SM}}\)](#).

Preliminary

Part III

Application Programming Interface

Preliminary

3 Data Types

3.1 Capability

A **Capability** (CAP) is a reference to a resource plus associated auxiliary data, such as access permissions.

Capabilities are opaque and immutable for applications – they cannot be inspected or modified directly; instead applications refer to a **Capability** via a **Capability Selector** (SEL).

3.1.1 Null Capability

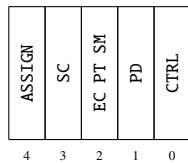
A **Null Capability** (CAP₀) does not refer to anything and carries no permissions.

3.1.2 Object Capability

An **Object Capability** (CAP_{OBJ}) is stored in the **Object Space** (SPC_{OBJ}) of a **PD** and refers to a kernel object.

3.1.2.1 PD Object Capability

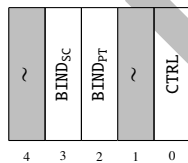
A **PD Object Capability** (CAP_{OBJ_{PD}}) refers to a **Protection Domain** (PD) and carries the following permissions:



CTRL [ctrl_pd](#) permitted if set.
 PD [create_pd](#) permitted if set.
 EC PT SM [create_ec](#), [create_pt](#), [create_sm](#) permitted if set.
 SC [create_sc](#) permitted if set.
 ASSIGN [assign_dev](#) permitted if set.

3.1.2.2 EC Object Capability


An **EC Object Capability** (CAP_{OBJ_{EC}}) refers to an **Execution Context** (EC) and carries the following permissions:



CTRL [ctrl_ec](#) permitted if set.
 BIND_{PT} [create_pt](#) can bind a **Portal** (PT) to the **EC** if set.
 BIND_{SC} [create_sc](#) can bind a **Scheduling Context** (SC) to the **EC** if set.

3.1.2.3 SC Object Capability

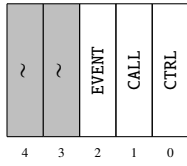
An **SC Object Capability** (CAP_{OBJ_{SC}}) refers to a **Scheduling Context** (SC) and carries the following permissions:



CTRL [ctrl_sc](#) permitted if set.

3.1.2.4 PT Object Capability

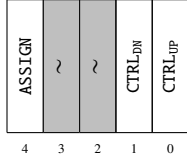
A **PT Object Capability** ($CAP_{OBJ_{PT}}$) refers to a **Portal (PT)** and carries the following permissions:



CTRL	$ctrl_pt$ permitted if set.
CALL	ipc_call permitted if set.
EVENT	Delivery of events permitted if set.

3.1.2.5 SM Object Capability

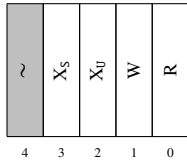
An **SM Object Capability** ($CAP_{OBJ_{SM}}$) refers to a **Semaphore (SM)** and carries the following permissions:



CTRL _{UP}	$ctrl_sm$ (Up) permitted if set.
CTRL _{DN}	$ctrl_sm$ (Down) permitted if set.
ASSIGN [†]	$assign_int$ permitted if set.

3.1.3 Memory Capability

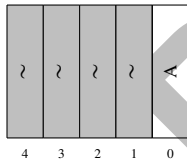
A **Memory Capability** (CAP_{MEM}) is stored in the **Memory Space (SPC_{MEM})** of a **PD**, refers to a 4 KiB page frame, and carries the following permissions:



R	the page frame is readable if set.
W	the page frame is writable if set.
X _U [‡]	the page frame is executable (in user mode) if set.
X _S [‡]	the page frame is executable (in supervisor mode) if set.

3.1.4 I/O Port Capability

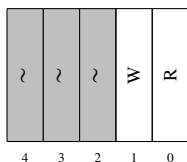
A **I/O Port Capability** (CAP_{PIO}) is stored in the **I/O Port Space (SPC_{PIO})** of a **PD**, refers to an I/O port, and carries the following permissions:



A	the I/O port is accessible (via IN/OUT) if set.
---	---

3.1.5 MSR Capability

A **MSR Capability** (CAP_{MSR}) is stored in the **MSR Space (SPC_{MSR})** of a **PD**, refers to a Model-Specific Register (**MSR**), and carries the following permissions:



R	the MSR is readable (via RDMSR) if set.
W	the MSR is writable (via WRMSR) if set.

[†]This permission bit is only defined for interrupt semaphores.

[‡]If the hardware supports only combined execute permissions (X) for both modes, then $X = X_U \vee X_S$.

3.2 Capability Selector

A **Capability Selector** (**SEL**) is an application-visible unsigned number as follows:

- An **Object Capability** Selector (**SEL_{OBJ}**) indexes into the **Object Space** (**SPC_{OBJ}**) of a **Protection Domain** (**PD**) and selects a slot that contains either a **Null Capability** (**CAP₀**) or an **Object Capability** (**CAP_{OBJ}**).
- A **Memory Capability** Selector (**SEL_{MEM}**) indexes into the **Memory Space** (**SPC_{MEM}**) of a **Protection Domain** (**PD**) and selects a slot that contains either a **Null Capability** (**CAP₀**) or a **Memory Capability** (**CAP_{MEM}**).
- An **I/O Port Capability** Selector (**SEL_{PIO}**) indexes into the **I/O Port Space** (**SPC_{PIO}**) of a **Protection Domain** (**PD**) and selects a slot that contains either a **Null Capability** (**CAP₀**) or an **I/O Port Capability** (**CAP_{PIO}**).
- An **MSR Capability** Selector (**SEL_{MSR}**) indexes into the **MSR Space** (**SPC_{MSR}**) of a **Protection Domain** (**PD**) and selects a slot that contains either a **Null Capability** (**CAP₀**) or an **MSR Capability** (**CAP_{MSR}**).

3.3 User Thread Control Block

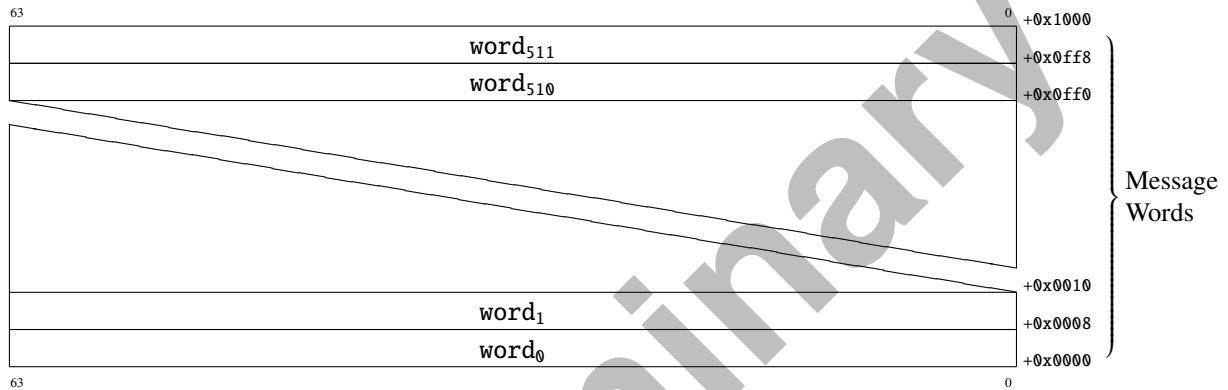
Each host **EC** (local/global thread) has its own **User Thread Control Block** (**ARM, x86**) (**UTCB**), which is mapped into the **Memory Space** (**SPC_{MEM}**) of the **PD** in which that **EC** is executing. A guest **EC** (virtual CPU) does not have a **UTCB**.

A **User Thread Control Block** (**ARM, x86**) has a size of one memory page (4 KiB). Because a **UTCB** is owned by the microhypervisor, it cannot be delegated using **ctrl_pd**.

To ensure proper visibility of loads and stores with relaxed memory ordering, application programs are expected to access a **UTCB** only from the **EC** to which that **UTCB** is bound.

3.3.1 Regular Layout

During regular **IPC** (see 3.4.1), the **UTCB** is used for data transfer and has a regular layout with 512 message words.



The data transfer from one **UTCB** to another **UTCB** is defined as follows:

- The data transfer is performed by the **CPU** on which the caller **EC** and callee **EC** execute.
- The data transfer uses the regular layout.
- The data is copied from low words to high words, beginning with word₀.
- The granularity of the loads and stores used for copying is **undefined**.
- Loads from and stores to the **UTCB** are **non-atomic** and use **relaxed** memory ordering.

3.3.2 Architectural Layout

During architectural **IPC** (see 3.4.2), the **UTCB** is used for state transfer and has an architectural layout (**ARM, x86**).

The state transfer between the architectural registers and a **UTCB** is defined as follows:

- The state transfer is performed by the **CPU** on which the faulting **EC** and callee **EC** execute.
- The state transfer uses the architectural layout.
- The state is copied between architectural registers and the **UTCB** in an **undefined** order.
- The granularity of the loads and stores used for copying is **undefined**.
- Loads from and stores to the **UTCB** are **non-atomic** and use **relaxed** memory ordering.

3.4 Message Transfer Descriptor

3.4.1 Regular IPC

For regular [Inter-Process Communication \(IPC\)](#), the [Message Transfer Descriptor \(ARM, x86\) \(MTD\)](#) is provided by the sender, passed to the receiver, and uses the following layout:

-	UTCB Message Words - 1
31	9 8 0

The [MTD](#) controls the data transfer (see [3.3.1](#)) as shown in [Figure 3.1](#):

- During [ipc_call](#), it specifies the number of message words to transfer from the [UTCB](#) of the caller [EC](#) (sender) to the [UTCB](#) of the callee [EC](#) (receiver).
- During [ipc_reply](#), it specifies the number of message words to transfer from the [UTCB](#) of the callee [EC](#) (sender) to the [UTCB](#) of the caller [EC](#) (receiver).

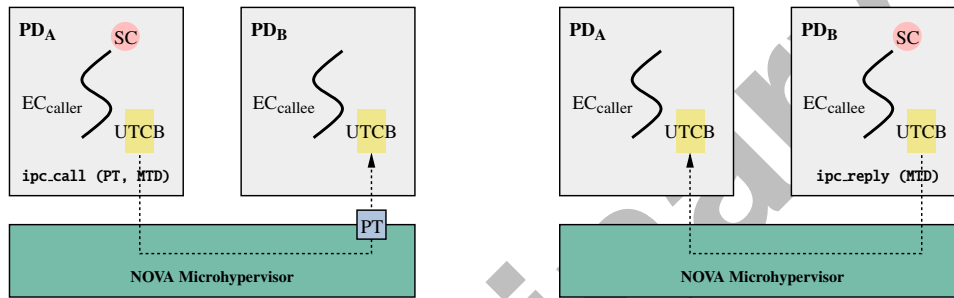


Figure 3.1: Regular [IPC](#)

3.4.2 Architectural IPC

For exceptions and intercepts, the [Message Transfer Descriptor \(ARM, x86\) \(MTD\)](#) is provided by the architectural event-specific portal ([ARM, x86](#)) or sender, passed to the receiver, and uses an architectural bitfield layout ([ARM, x86](#)):

- If a bit is 0, then the microhypervisor does **not** transmit the architectural state associated with that bit.
- If a bit is 1, then the microhypervisor transmits the architectural state associated with that bit.

The [MTD](#) controls the state transfer (see [3.3.2](#)) as shown in [Figure 3.2](#):

- During an exception/intercept, it specifies the subset of registers to transfer from the architectural state of the faulting [EC](#) (sender) to the [UTCB](#) of the callee [EC](#) (receiver).
- During [ipc_reply](#), it specifies the subset of registers to transfer from the [UTCB](#) of the callee [EC](#) (sender) to the architectural state of the faulting [EC](#) (receiver).

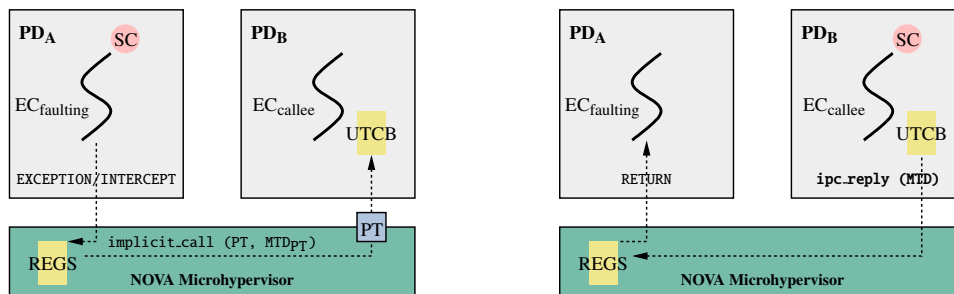


Figure 3.2: Architectural [IPC](#)

4 Hypercalls

4.1 Definitions

4.1.1 Hypercall Numbers

Each hypercall is identified by a unique number. The following hypercalls are currently defined:

Number	Hypercall	Section
0x0	ipc_call	4.2.1
0x1	ipc_reply	4.2.2
0x2	create_pd	4.3.1
0x3	create_ec	4.3.2
0x4	create_sc	4.3.3
0x5	create_pt	4.3.4
0x6	create_sm	4.3.5
0x7	ctrl_pd	4.4.1
0x8	ctrl_ec	4.4.2
0x9	ctrl_sc	4.4.3
0xa	ctrl_pt	4.4.4
0xb	ctrl_sm	4.4.5
0xc	ctrl_pm	4.5.1
0xd	assign_int	4.5.2
0xe	assign_dev	4.5.3
0xf	<i>reserved for future use</i>	

4.1.2 Status Codes

Hypercalls return a status code to indicate success or failure. The following status codes are currently defined:

Number	Status Code	Description
0x0	SUCCESS	Operation Successful
0x1	TIMEOUT	Operation Timeout
0x2	ABORTED	Operation Abort
0x3	OVRFLOW	Operation Overflow
0x4	BAD_HYP	Invalid Hypercall
0x5	BAD_CAP	Invalid Capability
0x6	BAD_PAR	Invalid Parameter
0x7	BAD_FTR	Invalid Feature
0x8	BAD_CPU	Invalid CPU Number
0x9	BAD_DEV	Invalid Device ID
0xa	INS_MEM	Insufficient Memory
≥0xb	<i>reserved for future use</i>	

4.1.3 Space Type

The following table lists the currently defined space types and for which architectures they are valid (✓):

Number	TYPE _{SPC}	ARM	x86	Description
0x0	SPC _{OBJ}	✓	✓	Object Space
0x1	SPC _{MEM}	✓	✓	Memory Space
0x2	SPC _{PIO}	×	✓	I/O Port Space
0x3	SPC _{MSR}	×	✓	MSR Space
≥0x4	<i>reserved for future use</i>			

4.1.4 Access Type

The following table lists the currently defined access types and for which space types they are valid (✓):

Number	TYPE _{ACC}	SPC _{OBJ}	SPC _{MEM}	SPC _{PIO}	SPC _{MSR}	Description
0x0	CPU_HST	✓	✓	✓	×	CPU Access from Host
0x1	CPU_GST	×	✓	✓	✓	CPU Access from Guest
0x2	DMA_HST	×	✓	×	×	DMA Access from Host
0x3	DMA_GST	×	✓	×	×	DMA Access from Guest
≥0x4	<i>reserved for future use</i>					

4.2 Communication

4.2.1 IPC Call

Parameters:

```
status = ipc_call (SEL_OBJ pt,          // Portal
                  MTD &mtd);          // Message Transfer Descriptor
```

Flags:

0	0	0	T
3	2	1	0

Description:

Sends a message from **EC_{CURRENT}** (caller) to the **EC** (callee) to which the specified **Portal (PT)** is bound.

Prior to the hypercall:

- { **PD_{CURRENT}**, **SEL_OBJ pt** } must refer to a **PT Object Capability (CAP_{OBJ_PT})** with permission **CALL**.

If the hypercall completed successfully:

- If **T=0 (No Timeout)**: If the callee **EC** was still busy handling a prior **ipc_call**, then the caller **EC** has helped run that prior **ipc_call** to completion, i.e. until the callee **EC** became available again.
- The microhypervisor has transferred a message from the **UTCB** of the caller **EC** to the **UTCB** of the callee **EC**. The content of that message is defined by the **MTD mtd**, which has been passed from the caller **EC** to the callee **EC**.
- The hypercall returns once the callee **EC** has issued an **ipc_reply**. Upon return, the **UTCB** of the caller **EC** and the parameter **mtd** have been updated by the reply message.
- The Current Scheduling Context (**SC_{CURRENT}**) has been donated to the callee **EC** upon **ipc_call** and returned back upon **ipc_reply**, thereby accounting the entire handling of the request to **SC_{CURRENT}**.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_CAP

- { **PD_{CURRENT}**, **SEL_OBJ pt** } did not refer to a **PT Object Capability (CAP_{OBJ_PT})** or that capability had insufficient permissions.

BAD_CPU

- Caller **EC** and callee **EC** are on different CPUs.

TIMEOUT

- The callee **EC** is still busy handling a prior **ipc_call** – only if **T=1 (Timeout)**.

ABORTED

- The callee **EC** is dead and the operation aborted.

4.2.2 IPC Reply

Parameters:

```
pid = ipc_reply (MTD &mtd);           // Message Transfer Descriptor
```

Flags:

0	0	0	0
3	2	1	0

Description:

Sends a reply message from **EC_{CURRENT}** (callee) back to the caller **EC** (if one exists) and subsequently waits for the next incoming message.

If the hypercall completed successfully:

- If a caller **EC** exists:
 - The microhypervisor has transferred a reply message from the **UTCB** of the callee **EC** back to the **UTCB** of the caller **EC**.
 - The content of that reply message is defined by the **MTD** **mtd**, which has been passed from the callee **EC** back to the caller **EC**.
 - The Current Scheduling Context (**SC_{CURRENT}**) that had been donated to the callee **EC** upon **ipc_call** has been returned back to the caller **EC**.
- **EC_{CURRENT}** blocks until the next incoming message arrives on any **Portal (PT)** bound to it.

Status:

This hypercall does not return directly.

Instead, when the next message arrives via a subsequent **ipc_call** to any **Portal (PT)** bound to the callee **EC**:

- The microhypervisor passes the Portal Identifier (**PID**) of the called **PT** to the callee **EC**.
- The **UTCB** of the callee **EC** and the parameter **mtd** have been updated by the incoming message.
- Execution of the callee **EC** continues at the Instruction Pointer (**IP**) configured in the called **PT**.

4.3 Object Creation

4.3.1 Create Protection Domain

Parameters:

```
status = create_pd (SEL_OBJ sel,          // Created PD
                   SEL_OBJ own);         // Owner PD
```

Flags:

0	0	0	0
3	2	1	0

Description:

Creates a new **Protection Domain (PD)**.

Prior to the hypercall:

- { **PD_{CURRENT}**, **SEL_{OBJ} own** } must refer to a **PD Object Capability (CAP_{OBJ_{PD}})** with permission PD.
- { **PD_{CURRENT}**, **SEL_{OBJ} sel** } must refer to a **Null Capability (CAP₀)**.

If the hypercall completed successfully:

- A new **Protection Domain (PD)** has been created.
- The resources for the created **PD** were accounted to the **PD** referred to by { **PD_{CURRENT}**, **SEL_{OBJ} own** }.
- { **PD_{CURRENT}**, **SEL_{OBJ} sel** } refers to a **PD Object Capability (CAP_{OBJ_{PD}})** for the created **PD** with **defined permissions** inherited from { **PD_{CURRENT}**, **SEL_{OBJ} own** }.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_CAP

- { **PD_{CURRENT}**, **SEL_{OBJ} own** } did not refer to a **PD Object Capability (CAP_{OBJ_{PD}})** or that capability had insufficient permissions.
- { **PD_{CURRENT}**, **SEL_{OBJ} sel** } did not refer to a **Null Capability (CAP₀)**.

INS_MEM

- { **PD_{CURRENT}**, **SEL_{OBJ} own** } had insufficient memory resources for **PD** creation.

4.3.2 Create Execution Context

Parameters:

```
status = create_ec (SEL_OBJ sel,          // Created EC
                   SEL_OBJ own,          // Owner PD
                   SEL_MEM utcb,         // UTCB Address (Page Number)
                   UINT cpu,             // CPU Number
                   UINT sp,              // Initial Stack Pointer
                   SEL_EVT evt);         // Event Selector Base
```

Flags:

0	F	V	T
3	2	1	0

Description:

Creates a new [Execution Context \(EC\)](#).

Prior to the hypercall:

- { [PD_{CURRENT}](#), [SEL_{OBJ} own](#) } must refer to a [PD Object Capability \(CAP_{OBJ_{PD}}\)](#) with permission EC.
- { [PD_{CURRENT}](#), [SEL_{OBJ} sel](#) } must refer to a [Null Capability \(CAP₀\)](#).

If the hypercall completed successfully:

- **V=0 (Thread):** A new host [Execution Context \(EC\)](#) has been created with its [UTCB](#) mapped at virtual page number [utcb](#) and its initial Stack Pointer ([SP](#)) set to [sp](#).
 - **T=0 (Local Thread):** [Portals \(PTs\)](#) can subsequently be bound to that [EC](#) and the [EC](#) will run whenever any of those bound portals is called.
 - **T=1 (Global Thread):** The [EC](#) will generate a startup exception the first time a [Scheduling Context \(SC\)](#) is bound to it.
- **V=1 (Virtual CPU):** A new guest [Execution Context \(EC\)](#) has been created. The [EC](#) will generate a startup exception the first time a [Scheduling Context \(SC\)](#) is bound to it. The parameters [utcb](#) and [sp](#) were ignored.
 - **T=0:** The virtual CPU uses no time adjustment.
 - **T=1:** The virtual CPU uses time offsetting.
- The created [EC](#) will be able to use [FPU](#) instructions only if the F-flag is set. Otherwise any [FPU](#) access by that [EC](#) will generate an exception.
- The created [EC](#) is bound to the [PD](#) referred to by { [PD_{CURRENT}](#), [SEL_{OBJ} own](#) } on [CPU](#) [cpu](#) with its Event Selector Base ([SEL_{EVT}](#)) set to [evt](#).
- The resources for the created [EC](#) were accounted to the [PD](#) referred to by { [PD_{CURRENT}](#), [SEL_{OBJ} own](#) }.
- { [PD_{CURRENT}](#), [SEL_{OBJ} sel](#) } refers to an [EC Object Capability \(CAP_{OBJ_{EC}}\)](#) for the created [EC](#) with all [defined permissions](#) set.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_CAP

- { [PD_{CURRENT}](#), [SEL_{OBJ} own](#) } did not refer to a [PD Object Capability \(CAP_{OBJ_{PD}}\)](#) or that capability had insufficient permissions.
- { [PD_{CURRENT}](#), [SEL_{OBJ} sel](#) } did not refer to a [Null Capability \(CAP₀\)](#).

BAD_CPU

- The CPU number is invalid.

BAD_FTR

- Virtual CPUs are not supported on the machine.

BAD_PAR

- UTCB region is not free or outside the user-accessible memory range.

INS_MEM

- { `PDCURRENT`, `SELOBJ own` } had insufficient memory resources for `EC` creation.

Preliminary

4.3.3 Create Scheduling Context

Parameters:

```
status = create_sc (SEL_OBJ sel,           // Created SC
                   SEL_OBJ own,           // Owner PD
                   SEL_OBJ ec,            // Bound EC
                   UINT budget,           // Scheduling Budget (in ms)
                   UINT priority);        // Scheduling Priority
```

Flags:

0	0	0	0
3	2	1	0

Description:

Creates a new [Scheduling Context \(SC\)](#).

Prior to the hypercall:

- { [PD_{CURRENT}](#), [SEL_{OBJ} own](#) } must refer to a [PD Object Capability \(CAP_{OBJ_{PD}}\)](#) with permission SC.
- { [PD_{CURRENT}](#), [SEL_{OBJ} ec](#) } must refer to an [EC Object Capability \(CAP_{OBJ_{EC}}\)](#) with permission BIND_{SC}.
- { [PD_{CURRENT}](#), [SEL_{OBJ} sel](#) } must refer to a [Null Capability \(CAP₀\)](#).

If the hypercall completed successfully:

- A new [Scheduling Context \(SC\)](#) has been created.
- The created [SC](#) is bound to the [EC](#) referred to by { [PD_{CURRENT}](#), [SEL_{OBJ} ec](#) } on the [CPU](#) of that [EC](#) with its scheduling parameters set to [budget](#) and [priority](#).
- The resources for the created [SC](#) were accounted to the [PD](#) referred to by { [PD_{CURRENT}](#), [SEL_{OBJ} own](#) }.
- { [PD_{CURRENT}](#), [SEL_{OBJ} sel](#) } refers to an [SC Object Capability \(CAP_{OBJ_{SC}}\)](#) for the created [SC](#) with all [defined permissions](#) set.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_CAP

- { [PD_{CURRENT}](#), [SEL_{OBJ} own](#) } did not refer to a [PD Object Capability \(CAP_{OBJ_{PD}}\)](#) or that capability had insufficient permissions.
- { [PD_{CURRENT}](#), [SEL_{OBJ} ec](#) } did not refer to a [EC Object Capability \(CAP_{OBJ_{EC}}\)](#) or that capability had insufficient permissions.
- { [PD_{CURRENT}](#), [SEL_{OBJ} sel](#) } did not refer to a [Null Capability \(CAP₀\)](#).
- Binding the [SC](#) to the [EC](#) failed, e.g. because the [EC](#) is a local [EC](#).

BAD_PAR

- Scheduling budget or priority was zero.

INS_MEM

- { [PD_{CURRENT}](#), [SEL_{OBJ} own](#) } had insufficient memory resources for [SC](#) creation.

4.3.4 Create Portal

Parameters:

```
status = create_pt (SEL_OBJ sel,      // Created PT
                   SEL_OBJ own,      // Owner PD
                   SEL_OBJ ec,       // Bound EC
                   UINT ip);         // Instruction Pointer
```

Flags:

0	0	0	0
3	2	1	0

Description:

Creates a new **Portal (PT)**.

Prior to the hypercall:

- { **PD_{CURRENT}**, **SEL_{OBJ} own** } must refer to a **PD Object Capability (CAP_{OBJ_{PD}})** with permission PT.
- { **PD_{CURRENT}**, **SEL_{OBJ} ec** } must refer to an **EC Object Capability (CAP_{OBJ_{EC}})** with permission BIND_{PT}.
- { **PD_{CURRENT}**, **SEL_{OBJ} sel** } must refer to a **Null Capability (CAP₀)**.

If the hypercall completed successfully:

- A new **Portal (PT)** has been created.
- The created **PT** is bound to the **EC** referred to by { **PD_{CURRENT}**, **SEL_{OBJ} ec** } on the **CPU** of that **EC**, with its portal Instruction Pointer (**IP**) set to **ip**, its initial **MTD** set to 0 and its initial **PID** set to 0.
- The resources for the created **PT** were accounted to the **PD** referred to by { **PD_{CURRENT}**, **SEL_{OBJ} own** }.
- { **PD_{CURRENT}**, **SEL_{OBJ} sel** } refers to an **PT Object Capability (CAP_{OBJ_{PT}})** for the created **PT** with all defined permissions set.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_CAP

- { **PD_{CURRENT}**, **SEL_{OBJ} own** } did not refer to a **PD Object Capability (CAP_{OBJ_{PD}})** or that capability had insufficient permissions.
- { **PD_{CURRENT}**, **SEL_{OBJ} ec** } did not refer to a **EC Object Capability (CAP_{OBJ_{EC}})** or that capability had insufficient permissions.
- { **PD_{CURRENT}**, **SEL_{OBJ} sel** } did not refer to a **Null Capability (CAP₀)**.
- Binding the **PT** to the **EC** failed, e.g. because the **EC** is not a local **EC**.

INS_MEM

- { **PD_{CURRENT}**, **SEL_{OBJ} own** } had insufficient memory resources for **PT** creation.

4.3.5 Create Semaphore

Parameters:

```
status = create_sm (SEL_OBJ sel,          // Created SM
                   SEL_OBJ own,          // Owner PD
                   UINT cnt);           // Initial Counter Value
```

Flags:

0	0	0	0
3	2	1	0

Description:

Creates a new **Semaphore (SM)**.

Prior to the hypercall:

- { **PD_{CURRENT}**, **SEL_OBJ own** } must refer to a **PD Object Capability (CAP_{OBJ_{PD}})** with permission SM.
- { **PD_{CURRENT}**, **SEL_OBJ sel** } must refer to a **Null Capability (CAP₀)**.

If the hypercall completed successfully:

- A new **Semaphore (SM)** has been created.
- The created **SM** has its initial counter value set to cnt.
- The resources for the created **SM** were accounted to the **PD** referred to by { **PD_{CURRENT}**, **SEL_OBJ own** }.
- { **PD_{CURRENT}**, **SEL_OBJ sel** } refers to an **SM Object Capability (CAP_{OBJ_{SM}})** for the created **SM** with all **defined permissions** set.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_CAP

- { **PD_{CURRENT}**, **SEL_OBJ own** } did not refer to a **PD Object Capability (CAP_{OBJ_{PD}})** or that capability had insufficient permissions.
- { **PD_{CURRENT}**, **SEL_OBJ sel** } did not refer to a **Null Capability (CAP₀)**.

INS_MEM

- { **PD_{CURRENT}**, **SEL_OBJ own** } had insufficient memory resources for **SM** creation.


4.4 Object Control

4.4.1 Control Protection Domain

Parameters:

```
status = ctrl_pd (SEL_OBJ spd,           // Protection Domain: Source
                  SEL_OBJ dpd,           // Protection Domain: Destination
                  SEL src,                // Base Selector: Source
                  SEL dst,                // Base Selector: Destination
                  UINT ord,               // Order
                  UINT pmm,               // Permission Mask
                  TYPE_SPC spc,           // Space Type
                  TYPE_ACC acc,           // Access Type
                  ATTR_CA ca,             // Cacheability Attribute
                  ATTR_SH sh);            // Shareability Attribute
```

Flags:



Description:

Takes capabilities from the Source **Protection Domain (PD)** and grants them to the Destination **Protection Domain (PD)** and thereby optionally reduces the permissions of the destination capabilities.

Prior to the hypercall:

- { **PD_{CURRENT}**, **SEL_{OBJ} spd** } must refer to a **PD Object Capability (CAP_{OBJ_{PD}})** with permission CTRL.
- { **PD_{CURRENT}**, **SEL_{OBJ} dpd** } must refer to a **PD Object Capability (CAP_{OBJ_{PD}})** with permission CTRL.
- { **PD_{CURRENT}**, **SEL_{OBJ} dpd** } must not refer to a **PD Object Capability (CAP_{OBJ_{PD}})** for **PD_{NOVA}**.
- **SEL src** and **SEL dst** must be order-aligned, i.e. $\text{src} \equiv 0 \pmod{2^{\text{ord}}}$ and $\text{dst} \equiv 0 \pmod{2^{\text{ord}}}$.
- **TYPE_{SPC} spc** and **TYPE_{ACC} acc** must be valid, i.e. supported by the architecture.
- **ATTR_{CA} ca** and **ATTR_{SH} sh** must be valid, i.e. supported by the architecture.

If the hypercall completed successfully:

- If **spc=SPC_{OBJ}**: All **CAP_{OBJ}** and **CAP₀** from source **SEL** range { **PD spd**, **SEL_{OBJ} src...src+2^{ord}-1** } were delegated to destination **SEL** range { **PD dpd**, **SEL_{OBJ} dst...dst+2^{ord}-1** }. Any pre-existing **CAP_{OBJ}** in the destination selector range were revoked. The parameters **acc**, **ca** and **sh** were ignored.
- If **spc=SPC_{MEM}**: All **CAP_{MEM}** and **CAP₀** from source **SEL** range { **PD spd**, **SEL_{MEM} src...src+2^{ord}-1** } were delegated to destination **SEL** range { **PD dpd**, **SEL_{MEM} dst...dst+2^{ord}-1** }. Any pre-existing **CAP_{MEM}** in the destination selector range were revoked.
- If **spc=SPC_{PIO}**: All **CAP_{PIO}** and **CAP₀** from source **SEL** range { **PD spd**, **SEL_{PIO} src...src+2^{ord}-1** } were delegated to destination **SEL** range { **PD dpd**, **SEL_{PIO} dst...dst+2^{ord}-1** }. Any pre-existing **CAP_{PIO}** in the destination selector range were revoked. The parameters **ca** and **sh** were ignored.
- If **spc=SPC_{MSR}**: All **CAP_{MSR}** and **CAP₀** from source **SEL** range { **PD spd**, **SEL_{MSR} src...src+2^{ord}-1** } were delegated to destination **SEL** range { **PD dpd**, **SEL_{MSR} dst...dst+2^{ord}-1** }. Any pre-existing **CAP_{MSR}** in the destination selector range were revoked. The parameters **ca** and **sh** were ignored.
- The permissions of each destination capability were masked by computing the logical AND of the permissions of the respective source capability and the permission mask **pmm**, i.e.
 - for bits set (1) in **pmm**, the respective permissions were *inherited* from the source capability.
 - for bits clear (0) in **pmm**, the respective permissions were *removed* for the destination capability.
- If the source capability was a **Null Capability (CAP₀)** or if the destination capability has zero permissions after masking, then the destination capability is now a **Null Capability (CAP₀)**.
- The resources for storing the granted capabilities were accounted to the **PD** referred to by { **PD_{CURRENT}**, **SEL_{OBJ} dpd** }.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_CAP

- { $PD_{CURRENT}$, SEL_{OBJ} spd } did not refer to a PD Object Capability ($CAP_{OBJ_{pd}}$) or that capability had insufficient permissions.
- { $PD_{CURRENT}$, SEL_{OBJ} dpd } did not refer to a PD Object Capability ($CAP_{OBJ_{pd}}$) or that capability had insufficient permissions.
- { $PD_{CURRENT}$, SEL_{OBJ} dpd } referred to a PD Object Capability ($CAP_{OBJ_{pd}}$) for PD_{NOVA} .

BAD_PAR

- SEL src or SEL dst was not order-aligned.
- SEL $src+2^{ord}-1$ or SEL $dst+2^{ord}-1$ was larger than the maximum selector number.
- If $spc=SPC_{PIO}$ or $spc=SPC_{MSR}$: SEL src was not equal to SEL dst .
- $TYPE_{SPC}$ spc or $TYPE_{ACC}$ acc was not valid, i.e. not supported by the architecture.
- $ATTR_{CA}$ ca or $ATTR_{SH}$ sh was not valid, i.e. not supported by the architecture.

INS_MEM

- { $PD_{CURRENT}$, SEL_{OBJ} dpd } had insufficient memory resources for allocating the storage required for granting all destination capabilities. This constitutes a partial failure of the operation, because all destination capabilities up to the first allocation failure have been granted.

4.4.2 Control Execution Context

Parameters:

```
status = ctrl_ec (SEL_OBJ ec);           // Execution Context
```

Flags:

0	0	0	S
3	2	1	0

Description:

Prior to the hypercall:

- { $PD_{CURRENT}$, SEL_{OBJ} ec } must refer to a **EC Object Capability** ($CAP_{OBJ_{EC}}$) with permission CTRL.

If the hypercall completed successfully:

- The **EC** referred to by { $PD_{CURRENT}$, SEL_{OBJ} ec } has been forced to enter the microhypervisor. It will generate a recall exception prior to its next exit from the microhypervisor and will traverse through the respective **Portal** (PT).
- If **S=0 (Weak)**: the hypercall returns as soon as the recall exception has been *pending*, i.e. the EC may not have entered the microhypervisor yet.
- If **S=1 (Strong)**: the hypercall returns as soon as the recall exception has been *observed*, i.e the EC will have entered the microhypervisor.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_CAP

- { $PD_{CURRENT}$, SEL_{OBJ} ec } did not refer to a **EC Object Capability** ($CAP_{OBJ_{EC}}$) or that capability had insufficient permissions.

4.4.3 Control Scheduling Context

Parameters:

```
status = ctrl_sc (SEL_OBJ sc,          // Scheduling Context
                  UINT &ticks);       // Total Consumed Execution Time
```

Flags:

0	0	0	0
3	2	1	0

Description:

Prior to the hypercall:

- { `PDCURRENT`, `SELOBJ sc` } must refer to an **SC Object Capability** (`CAPOBJsc`) with permission CTRL.

If the hypercall completed successfully:

- The microhypervisor has returned the total consumed execution time in `ticks` for the **SC** referred to by { `PDCURRENT`, `SELOBJ sc` }.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_CAP

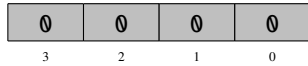
- { `PDCURRENT`, `SELOBJ sc` } did not refer to an **SC Object Capability** (`CAPOBJsc`) or that capability had insufficient permissions.

4.4.4 Control Portal

Parameters:

```
status = ctrl.pt (SEL_OBJ pt,          // Portal
                  UINT  pid,          // Portal Identifier
                  MTD   mtd);         // Message Transfer Descriptor
```

Flags:



Description:

Prior to the hypercall:

- { `PDCURRENT`, `SELOBJ pt` } must refer to a **PT Object Capability** (`CAPOBJ_PT`) with permission CTRL.

If the hypercall completed successfully:

- The microhypervisor has set the Portal Identifier (**PID**) to `pid` and the **Message Transfer Descriptor** (**ARM, x86**) (**MTD**) to `mtd` for the **Portal** referred to by { `PDCURRENT`, `SELOBJ pt` }.
- Subsequent portal traversals will use the new **MTD** and return the new **PID**.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_CAP

- { `PDCURRENT`, `SELOBJ pt` } did not refer to a **PT Object Capability** (`CAPOBJ_PT`) or that capability had insufficient permissions.

4.4.5 Control Semaphore

Parameters:

```
status = ctrl_sm (SEL_OBJ sm,          // Semaphore
                  UINT ticks);         // Deadline Timeout
```

Flags:

0	0	Z	D
3	2	1	0

Description:

Prior to the hypercall:

- If **D=0 (Up)**: { **PD_{CURRENT}**, **SEL_{OBJ} sm** } must refer to a **SM Object Capability (CAP_{OBJ_{SM}})** with permission **CTRL_{UP}**.
- If **D=1 (Down)**: { **PD_{CURRENT}**, **SEL_{OBJ} sm** } must refer to a **SM Object Capability (CAP_{OBJ_{SM}})** with permission **CTRL_{DN}**.

If the hypercall completed successfully:

- If **D=0 (Up)**: if there were **ECs** blocked on the semaphore, then the microhypervisor has released one of those blocked **ECs**. Otherwise, the microhypervisor has incremented the semaphore counter. The deadline timeout value and the Z-flag were ignored.
- If **D=1 (Down)**: if the semaphore counter was larger than zero, then the microhypervisor has decremented the semaphore counter (**Z=0**) or set it to zero (**Z=1**). Otherwise, the microhypervisor has blocked **EC_{CURRENT}** on the semaphore. If the deadline timeout value was non-zero, **EC_{CURRENT}** unblocks with a timeout status when the architectural timer reaches or exceeds the specified ticks value.

Blocking and releasing of **ECs** on a semaphore uses the FIFO queueing discipline.

Status:

SUCCESS

- The hypercall completed successfully.

TIMEOUT

- If **D=1**: Down operation aborted when the timeout triggered.

OVRFLOW

- If **D=0**: Up operation aborted because the semaphore counter would overflow.

BAD_CAP

- { **PD_{CURRENT}**, **SEL_{OBJ} sm** } did not refer to a **SM Object Capability (CAP_{OBJ_{SM}})** or that capability had insufficient permissions.

BAD_CPU

- If **D=1** on an interrupt semaphore: Attempt to wait for the interrupt on a different **CPU** than the **CPU** to which that interrupt has been routed via **assign.int**.


4.5 Platform Management

4.5.1 Control Power Management

Parameters:

```
status = ctrl_pm (UINT state);           // State Information
```

Flags:



Description:

Transitions the platform to the specified power management state.

Prior to the hypercall:

- $PD_{CURRENT}$ must be the [Root Protection Domain](#) (PD_{ROOT}).
- If $OP=1$ (S-State Transition):

- The state parameter uses the following encoding:



- The value S designates the state to enter. The values A and B are the first two bytes of the respective $_Sx$ package in the [ACPI](#) root namespace as follows:

S	A	B	Returns	Description
0x1	_S1[0]	_S1[1]	✓	S1: Power-On Suspend
0x2	_S2[0]	_S2[1]	✓	S2: Standby
0x3	_S3[0]	_S3[1]	✓	S3: Suspend to RAM
0x4	_S4[0]	_S4[1]	×	S4: Suspend to Disk
0x5	_S5[0]	_S5[1]	×	S5: Soft Off
0x7	0x0	0x0	×	Platform Reset

- The caller is responsible for invoking the necessary pre-sleep [ACPI](#) methods, for transitioning platform devices into a suitable Dx sleep state, and for programming wakeup events.

If the hypercall completed successfully:

- If $OP=1$ (S-State Transition): The platform enters the specified [ACPI](#) sleep state or resets.
 - For shallow sleep states, the hypercall returns upon a wakeup event. The caller is responsible for invoking the necessary post-sleep [ACPI](#) methods and for transitioning platform devices back into the D0 working state.
 - For deep sleep states or platform reset, the hypercall does not return.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_HYP

- The hypercall was not issued from the [Root Protection Domain](#) (PD_{ROOT}).

BAD_PAR

- The requested operation (OP) is not supported.

BAD_FTR

- The requested power management state is not supported.

ABORTED

- A concurrent power management request prevailed.

4.5.2 Assign Interrupt

Parameters:

```
status = assign_int (SEL_OBJ sm,           // Interrupt Semaphore
                    UINT  cpu,             // CPU Number
                    UINT  dev,             // MSI Authorized Device
                    UINT  &msi_addr,       // MSI Message Address
                    UINT  &msi_data);      // MSI Message Data
```

Flags:

G	P	T	M
3	2	1	0

Description:

Configures an interrupt and routes it to the specified CPU.

Prior to the hypercall:

- { `PDCURRENT`, `SEL_OBJ sm` } must refer to a `SM Object Capability (CAPOBJSM)` with permission `ASSIGN`.
- `CAPOBJSM` must refer to an interrupt semaphore and thereby designates the interrupt.

If the hypercall completed successfully:

- The interrupt referred to by { `PDCURRENT`, `SEL_OBJ sm` } has been routed to the CPU `cpu`.
- Mask
 - `M=0`: The interrupt is now unmasked, i.e. it will be signaled on the semaphore.
 - `M=1`: The interrupt is now masked, i.e. it will not be signaled on the semaphore.
- Trigger
 - `T=0`: The interrupt is now configured for edge-triggered operation.
 - `T=1`: The interrupt is now configured for level-triggered operation.
- Polarity
 - `P=0`: The interrupt is now configured for active-high operation.
 - `P=1`: The interrupt is now configured for active-low operation.
- Guest
 - `G=0`: The interrupt is now host-owned.
 - `G=1`: The interrupt is now guest-owned (VM pass-through).
- If the interrupt is an `MSI`, only the `PCI` device referred to by `dev` will be authorized to generate that `MSI`. The device driver must program the returned `msi_addr` and `msi_data` values into the `MSI` registers of that device to ensure proper interrupt operation. If the interrupt is pin-based, the parameter `dev` was ignored and the parameters `msi_addr` and `msi_data` return 0.

Prior to the first invocation of `assign_int` for an interrupt, the state of that interrupt is as follows:

- the interrupt is masked.
- trigger, polarity and ownership are undefined.
- target CPU and authorized device are undefined.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_CPU

- The specified CPU number was invalid.

BAD_CAP

- { $PD_{CURRENT}$, $SEL_{OBJ\ sm}$ } did not refer to a SM Object Capability ($CAP_{OBJ_{SM}}$) or that capability had insufficient permissions.
- $CAP_{OBJ_{SM}}$ did not refer to an interrupt semaphore.

Preliminary

4.5.3 Assign Device

Parameters:

```
status = assign_dev (SEL_OBJ pd,           // Protection Domain
                    SEL_MEM smmu,         // SMMU Address (Page Number)
                    UINT dev,             // Assigned Device (SID/BDF)
                    TYPE_ACC acc);        // Access Type
```

Flags:

0	0	0	0
3	2	1	0

Description:

Assigns the specified device (*) to the specified **Protection Domain (PD)**:

- **ARM**: dev encodes the **SID** of the device and also the **SMMU** resources (stream mapping group, translation context) to be used for managing that device.
- **x86**: dev encodes the **BDF** of the device. There are no **SMMU** resources needed.

Prior to the hypercall:

- **PD_{CURRENT}** must be the **Root Protection Domain (PD_{ROOT})**.
- { **PD_{CURRENT}**, **SEL_{OBJ} pd** } must refer to a **PD Object Capability (CAP_{OBJ_{pd}})** with permission **ASSIGN**.
- { **PD_{NOVA}**, **SEL_{MEM} smmu** } must refer to the physical address of an **SMMU** device.
- The **SID/BDF** and **SMMU** resources encoded in dev must be supported by the hardware (see 6.6.1).
- **TYPE_{ACC} acc** must refer to a DMA access type.

If the hypercall completed successfully:

- The device, referred to by the **SID/BDF** in dev, has been assigned to the **Protection Domain (PD)** referred to by { **PD_{CURRENT}**, **SEL_{OBJ} pd** }, such that DMA transactions of that device will be translated by the DMA page table corresponding to acc of that **PD**.
- DMA transactions of that device will be managed using the **SMMU** resources encoded in dev. Prior users of those **SMMU** resources have been unconfigured.

Status:

SUCCESS

- The hypercall completed successfully.

BAD_HYP

- The hypercall was not issued from the **Root Protection Domain (PD_{ROOT})**.

BAD_DEV

- { **PD_{NOVA}**, **SEL_{MEM} smmu** } did not refer to the physical address of an **SMMU** device.

BAD_CAP

- { **PD_{CURRENT}**, **SEL_{OBJ} pd** } did not refer to a **PD Object Capability (CAP_{OBJ_{pd}})** or that capability had insufficient permissions.

BAD_PAR

- At least one of the parameters dev or acc was not valid.

*See the architecture-specific binding for encoding details.

5 Booting

5.1 Microhypervisor

5.1.1 ELF Image Loading

The bootloader must place all loadable (PT_LOAD) program segments of the NOVA microhypervisor into physical memory (RAM) according to the physical addresses (`p_paddr`) and memory sizes (`p_memsz`) defined in the NOVA microhypervisor [ELF](#) image. The following is an example:

```
readelf -l nova.elf
```

Elf file type is EXEC (Executable file)

Entry point 0x48000000

There are 2 program headers, starting at offset 64

Program Headers:

Type	Offset	VirtAddr	PhysAddr
	FileSiz	MemSiz	Flags Align
LOAD	0x00000000000000b0	0x0000000048000000	0x0000000048000000
	0x0000000000000268	0x0000000000001000	RWE 0x8
LOAD	0x0000000000000800	0x0000ff8000001000	0x0000000048001000
	0x000000000000e960	0x000000000000fff000	RWE 0x800

If the physical address range defined in the [ELF](#) image is suboptimal for a particular platform, the bootloader may optionally shift all loadable program segments lower or higher in physical memory, by applying an offset, subject to the following constraints:

- The same offset must be applied to each loadable program segment and to the entry point.
- The offset must be a multiple of 2 MiB, i.e. $\text{PhysAddr}_{\text{NEW}} = \text{PhysAddr}_{\text{ELF}} \pm n \times 2 \text{ MiB}$.
- The entire physical memory region occupied by the NOVA microhypervisor must be RAM.

After loading the NOVA microhypervisor into physical memory, the bootloader must invoke the entry point of the [ELF](#) image with architecture-specific preconditions ([ARM](#), [x86](#)).

5.1.2 Platform Resource Access

Possession of a [PD Object Capability](#) ($\text{CAP}_{\text{OBJ}_{\text{PD}}}$) for PD_{NOVA} allows the caller to invoke the `ctrl_pd` hypercall to take resources from the [NOVA Protection Domain](#) and grant them to another [Protection Domain](#).

The following capabilities can be taken from the [NOVA Protection Domain](#) (PD_{NOVA}):

Physical Memory

{ PD_{NOVA} , SEL_{MEM} $0 \dots \text{PHYS}_{\text{NUM}} - 1$ } refer to CAP_{MEM} for page frames in physical memory, where PHYS_{NUM} is the number of page frames supported by the platform. Physical memory regions protected by the NOVA microhypervisor ([ARM](#), [x86](#)) cannot be taken.

Interrupt Semaphores

{ PD_{NOVA} , SEL_{OBJ} $1024 \dots 1024 + \text{INT}_{\text{NUM}} - 1$ } refer to $\text{CAP}_{\text{OBJ}_{\text{SM}}}$ for interrupt semaphores, where INT_{NUM} is the number of supported interrupts, as conveyed by the [HIP](#). These capabilities can be used with the `ctrl_sm` and `assign_int` hypercalls.

Console Signaling Semaphore

{ PD_{NOVA} , SEL_{OBJ} $\text{SEL}_{\text{NUM}} - 1$ } refers to a $\text{CAP}_{\text{OBJ}_{\text{SM}}}$ for the signaling semaphore of the NOVA memory-buffer console. This capability can be used with the `ctrl_sm` hypercall.

5.2 Root Protection Domain

After the NOVA microhypervisor has initialized the system, it creates the following initial kernel objects:

- PD_{ROOT} – the [Root Protection Domain](#)
- EC_{ROOT} – the [Root Execution Context](#) (executing in PD_{ROOT})
- SC_{ROOT} – the [Root Scheduling Context](#) (bound to EC_{ROOT})

The [Root Protection Domain](#) is responsible for bootstrapping the other components of the user-mode framework by creating additional kernel objects, loading additional images, assigning resources, etc.

5.2.1 ELF Image Format

The [ELF](#) image of the [Root Protection Domain](#) (PD_{ROOT}) must be an executable (ET_EXEC) file that has been compiled for the respective architecture and

- linked such that $p_filesz = p_memsz$
- loaded such that $p_vaddr \equiv LOAD_ADDR^* + p_offset \pmod{PAGE_SIZE}$

holds for each loadable (PT_LOAD) program segment. These constraints ensure that the NOVA microhypervisor can map all program segments directly from physical into virtual memory without any additional memory allocation or copying. The following is an example:

```
readelf -l root.elf
```

```
Elf file type is EXEC (Executable file)
```

```
Entry point 0x10000120
```

```
There are 2 program headers, starting at offset 64
```

```
Program Headers:
```

Type	Offset	VirtAddr	PhysAddr
	FileSiz	MemSiz	Flags Align
LOAD	0x0000000000000000	0x0000000001000000	0x0000000001000000
	0x0000000000000a75	0x0000000000000a75	R E 0x1000
LOAD	0x0000000000001000	0x0000000001000100	0x0000000001000100
	0x000000000000f004	0x000000000000f004	RW 0x1000

5.2.2 Initial Configuration

Prior to invoking the entry point of the [Root Protection Domain](#) (PD_{ROOT}) [ELF](#) image, using the [Root Execution Context](#) (EC_{ROOT}), the NOVA microhypervisor sets up PD_{ROOT} as follows.

5.2.2.1 Object Space

The object space contains the following initial capabilities:

- { PD_{ROOT} , SEL_{OBJ} $SEL_{NUM}-1$ } refers to a [PD Object Capability](#) ($CAP_{OBJ_{PD}}$) for PD_{NOVA} .
- { PD_{ROOT} , SEL_{OBJ} $SEL_{NUM}-2$ } refers to a [PD Object Capability](#) ($CAP_{OBJ_{PD}}$) for PD_{ROOT} .
- { PD_{ROOT} , SEL_{OBJ} $SEL_{NUM}-3$ } refers to a [EC Object Capability](#) ($CAP_{OBJ_{EC}}$) for EC_{ROOT} .
- { PD_{ROOT} , SEL_{OBJ} $SEL_{NUM}-4$ } refers to a [SC Object Capability](#) ($CAP_{OBJ_{SC}}$) for SC_{ROOT} .

All other { PD_{ROOT} , SEL_{OBJ} } refer to a [Null Capability](#) (CAP_0).

The value of SEL_{NUM} is conveyed in the [Hypervisor Information Page](#) (ARM, x86).

*This is the address in physical memory at which the bootloader has placed the ELF image.

5.2.2.2 Memory Space

ELF Program Segments

The microhypervisor maps the [Root Protection Domain](#) (PD_{ROOT}) into virtual memory according to the virtual addresses (p_vaddr), memory sizes (p_memsz) and page attributes (p_flags) of all loadable (PT_LOAD) program segments defined in the PD_{ROOT} ELF image.

Hypervisor Information Page

The microhypervisor maps the [Hypervisor Information Page](#) (ARM, x86) read-only into the memory space 4 KiB below the end of user-accessible virtual memory. The virtual address of the [HIP](#) is passed to EC_{ROOT} at the entry point (ARM, x86).

UTCB

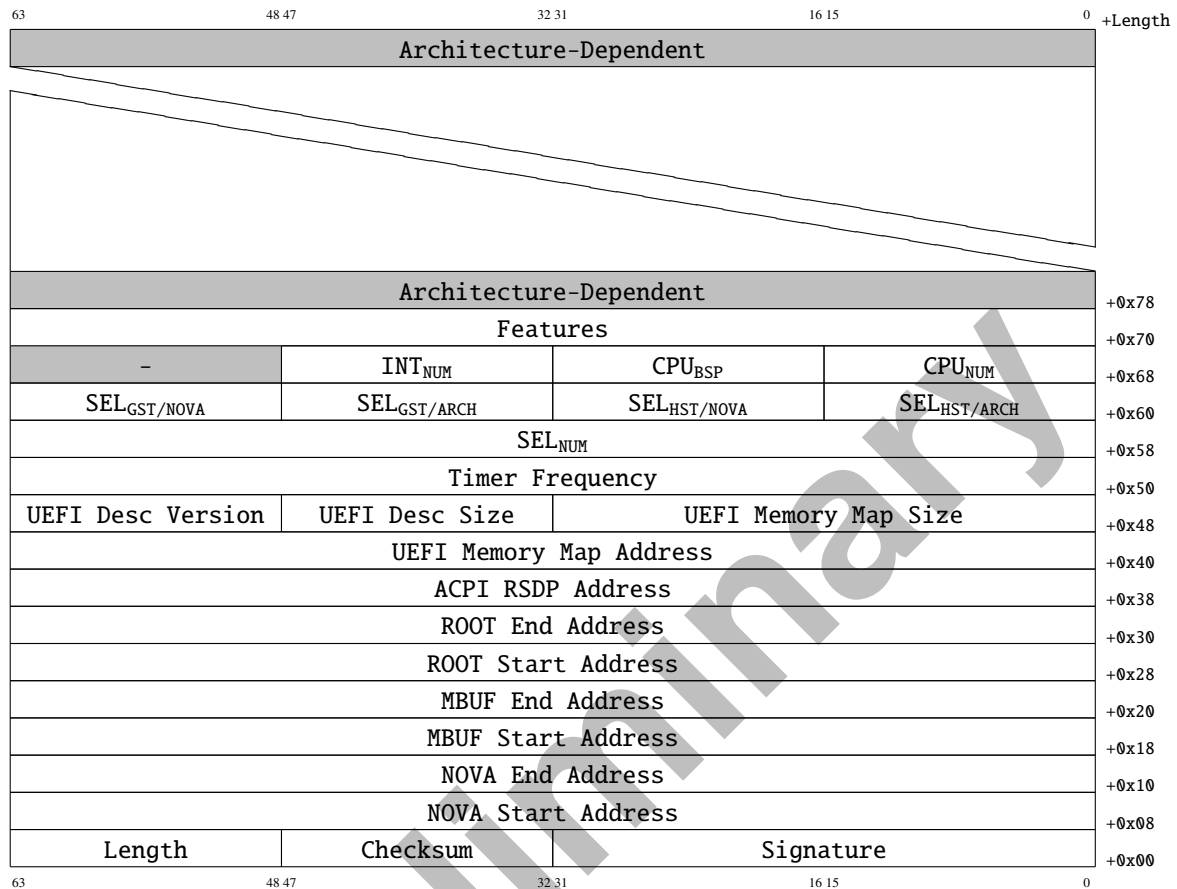
The microhypervisor maps the [User Thread Control Block](#) (ARM, x86) of EC_{ROOT} into the memory space 4 KiB below the address of the [Hypervisor Information Page](#) (ARM, x86).

All other { PD_{ROOT} , SEL_{MEM} } refer to a [Null Capability](#) (CAP_0).

Preliminary

5.3 Hypervisor Information Page

The [Hypervisor Information Page \(ARM, x86\) \(HIP\)](#) conveys information about the platform and configuration to the [Root Protection Domain \(PD_{ROOT}\)](#) and has the following layout:



All HIP fields are unsigned values, unless stated otherwise, and have the following meaning:

Signature

The value `0x41564f4e` identifies the NOVA microhypervisor.

Checksum

The checksum is valid if 16bit-wise addition of the entire [HIP](#) contents produces a value of `0`.

Length

Length of the entire [HIP](#) in bytes.

NOVA Start/End Address

Physical start and end address of the NOVA microhypervisor image.

MBUF Start/End Address

Physical start and end address of the memory buffer console region (see [C.1](#)).

ROOT Start/End Address

Physical start and end address of the root protection domain image.

ACPI RSDP Address

Physical address of the [ACPI](#) Root System Description Pointer (`0xffffffffffffffff` if not present).

UEFI Memory Map Address

Physical address of the [UEFI](#) Memory Map (`0xffffffffffffffff` if not present).

UEFI Memory Map Size

Total size of the [UEFI](#) Memory Map (0 if not present).

UEFI Desc Size

[UEFI](#) Memory Descriptor Size (0 if not present).

UEFI Desc Version

[UEFI](#) Memory Descriptor Version (0 if not present).

Timer Frequency

Timer tick frequency in Hz.

SEL_{NUM}

Total number of capability selectors in each object space.

SEL_{HST/ARCH}

Number of capability selectors required for handling architectural host events. ([ARM](#), [x86](#))

SEL_{HST/NOVA}

Number of additional capability selectors required for handling microhypervisor host events. ([ARM](#), [x86](#))

SEL_{GST/ARCH}

Number of capability selectors required for handling architectural guest events. ([ARM](#), [x86](#))

SEL_{GST/NOVA}

Number of additional capability selectors required for handling microhypervisor guest events. ([ARM](#), [x86](#))

CPU_{NUM}

Total number of CPUs that are online.

CPU_{BSP}

The bootstrap CPU on which [EC_{ROOT}](#) and [SC_{ROOT}](#) have been created.

INT_{NUM}

Total number of interrupts that can be used via interrupt semaphores.

Features

Supported platform features.

Architecture-Dependent

Architecture-dependent part. ([ARM](#), [x86](#))

Part IV

Application Binary Interface

6 ABI aarch64

6.1 Boot State

6.1.1 NOVA Microhypervisor

The bootloader must set up the [CPU](#) register state according to one of the launch types listed below when it transfers control to the NOVA microhypervisor. Furthermore, the following preconditions must be satisfied:

- The [CPU](#) must execute in EL2 (hypervisor mode) or in EL3 (monitor mode).
- Paging (MMU) must be disabled (SCTLR.ELx.M=0) or must use an identity (1:1) mapping.
- Interrupts must be disabled (PSTATE.DAIF=0b1111).
- The physical memory region occupied by the microhypervisor image must be clean to the PoC.
- All [DMA](#) activity targeting the physical memory region occupied by the microhypervisor must be quiesced. That physical memory region should also be protected against [DMA](#) accesses on systems with an SMMU.

6.1.1.1 Multiboot v2 Launch

Only this launch type supports 64-bit [UEFI](#) platforms.

Register	Value / Description
IP	Physical address of the NOVA Protection Domain (PD_{NOVA}) ELF image entry point
X0	Multiboot v2 magic value (0x36d76289) [7]
X1	Physical address of the Multiboot v2 information structure [7]
Other	~

The NOVA microhypervisor consumes the following multiboot tags, if present: 1, 3, 12, 20.

6.1.1.2 Multiboot v1 Launch

Register	Value / Description
IP	Physical address of the NOVA Protection Domain (PD_{NOVA}) ELF image entry point
X0	Multiboot v1 magic value (0x2badb002) [6]
X1	Physical address of the Multiboot v1 information structure [6]
Other	~

The NOVA microhypervisor consumes the following multiboot flags, if present: 2, 3.

6.1.1.3 Legacy Launch

Register	Value / Description
IP	Physical address of the NOVA Protection Domain (PD_{NOVA}) ELF image entry point
X0	Physical address of the Flattened Device Tree [5] (FDT) for the system [†]
X1	Physical address of the Root Protection Domain (PD_{ROOT}) ELF image
Other	~

[†] Due to its alignment constraint, a valid FDT address will never be equal to a Multiboot magic value.

6.1.2 Root Protection Domain

The NOVA microhypervisor sets up the CPU register state as follows when it transfers control to the Root Execution Context (EC_{ROOT}):

Register	Value / Description
IP	Virtual address of the Root Protection Domain (PD_{ROOT}) ELF image entry point
SP	Virtual address of the Hypervisor Information Page (ARM, x86) (HIP)
X0	X0 at boot time [†]
X1	X1 at boot time [†]
X2	X2 at boot time [†]
Other	~

[†]The register contains the preserved original value from the point when control was transferred from the bootloader to the microhypervisor.

6.2 Protected Resources

The following resources are protected by the NOVA microhypervisor and are therefore **inaccessible** to user-mode applications.

6.2.1 Memory Space

- Physical memory occupied by the NOVA microhypervisor (conveyed via **HIP**).
- Physical memory occupied by GICD, GICR, GICC, GICH devices (conveyed via **ACPI MADT** or via **FDT**).
- Physical memory occupied by SMMU devices (conveyed via **ACPI IORT** or via **FDT**).
- Physical memory occupied by firmware runtime services (conveyed via **UEFI** memory map).

6.3 Physical Memory

6.3.1 Memory Map

The **Root Protection Domain (PD_{ROOT})** can obtain a list of available/reserved memory regions as follows:

- On platforms using Unified Extensible Firmware Interface [11], by parsing the **UEFI memory map**.
- On platforms using Flattened Device Tree [5], by parsing the **FDT**.

6.4 Virtual Memory

The accessible virtual memory range for user-mode applications is **0 – 0x7fffffffff**.

6.4.1 Cacheability Attributes

Encoding	ATTR_{CA}	Description
0x0	DEV	Device
0x1	DEV_E	Device, Early Ack
0x2	DEV_RE	Device, Early Ack, Reordering
0x3	DEV_GRE	Device, Early Ack, Reordering, Gathering
0x4	–	<i>reserved</i>
0x5	MEM_NC	Memory, Inner/Outer Non-Cacheable
0x6	MEM_WT	Memory, Inner/Outer Write-Through
0x7	MEM_WB	Memory, Inner/Outer Write-Back

Please refer to [3] for details on the architectural behavior.

6.4.2 Shareability Attributes

Encoding	ATTR_{SH}	Description
0x0	NONE	Not Shareable
0x1	–	<i>reserved</i>
0x2	OUTER	Outer Shareable
0x3	INNER	Inner Shareable

Please refer to [3] for details on the architectural behavior.

6.5 Event-Specific Capability Selectors

For the delivery of exception/intercept messages, the microhypervisor performs an implicit portal traversal.

The selector for the destination portal (SEL_{OBJ}):

- is determined by adding the exception/intercept number to the affected EC's Event Selector Base (SEL_{EVT}).
- indexes into the Object Space (SPC_{OBJ}) of the affected EC's Protection Domain (PD).
- must refer to a PT Object Capability ($CAP_{OBJ_{PT}}$), otherwise the affected EC is killed.

6.5.1 Architectural Events

Host Exceptions and Guest Intercepts

SEL_{OBJ}	Exception / Intercept	SEL_{OBJ}	Exception / Intercept
$SEL_{EVT} + 0x0$	Unknown Reason	$SEL_{EVT} + 0x20$	Instruction Abort (lower EL)
$SEL_{EVT} + 0x1$	Trapped WFI or WFE	$SEL_{EVT} + 0x21$	Instruction Abort (same EL)*
$SEL_{EVT} + 0x2$	reserved	$SEL_{EVT} + 0x22$	PC Alignment Fault
$SEL_{EVT} + 0x3$	Trapped MCR or MRC	$SEL_{EVT} + 0x23$	reserved
$SEL_{EVT} + 0x4$	Trapped MCRR or MRRC	$SEL_{EVT} + 0x24$	Data Abort (lower EL)
$SEL_{EVT} + 0x5$	Trapped MCR or MRC	$SEL_{EVT} + 0x25$	Data Abort (same EL)*
$SEL_{EVT} + 0x6$	Trapped LDC or STC	$SEL_{EVT} + 0x26$	SP Alignment Fault
$SEL_{EVT} + 0x7$	SVE, SIMD, FPU	$SEL_{EVT} + 0x27$	reserved
$SEL_{EVT} + 0x8$	Trapped VMRS Access	$SEL_{EVT} + 0x28$	Trapped FPU (AArch32)
$SEL_{EVT} + 0x9$	Trapped PAAuth Instruction	$SEL_{EVT} + 0x29$	reserved
$SEL_{EVT} + 0xa$	reserved	$SEL_{EVT} + 0x2a$	reserved
$SEL_{EVT} + 0xb$	reserved	$SEL_{EVT} + 0x2b$	reserved
$SEL_{EVT} + 0xc$	Trapped MRRC	$SEL_{EVT} + 0x2c$	Trapped FPU (AArch64)
$SEL_{EVT} + 0xd$	Branch Target Exception	$SEL_{EVT} + 0x2d$	reserved
$SEL_{EVT} + 0xe$	Illegal Execution State	$SEL_{EVT} + 0x2e$	reserved
$SEL_{EVT} + 0xf$	reserved	$SEL_{EVT} + 0x2f$	Error
$SEL_{EVT} + 0x10$	reserved	$SEL_{EVT} + 0x30$	Breakpoint (lower EL)
$SEL_{EVT} + 0x11$	SVC (from AArch32 State)	$SEL_{EVT} + 0x31$	Breakpoint (same EL)*
$SEL_{EVT} + 0x12$	HVC (from AArch32 State)	$SEL_{EVT} + 0x32$	Software Step (lower EL)
$SEL_{EVT} + 0x13$	SMC (from AArch32 State)	$SEL_{EVT} + 0x33$	Software Step (same EL)*
$SEL_{EVT} + 0x14$	reserved	$SEL_{EVT} + 0x34$	Watchpoint (lower EL)
$SEL_{EVT} + 0x15$	SVC (from AArch64 State)*	$SEL_{EVT} + 0x35$	Watchpoint (same EL)*
$SEL_{EVT} + 0x16$	HVC (from AArch64 State)	$SEL_{EVT} + 0x36$	reserved
$SEL_{EVT} + 0x17$	SMC (from AArch64 State)	$SEL_{EVT} + 0x37$	reserved
$SEL_{EVT} + 0x18$	Trapped MSR or MRS	$SEL_{EVT} + 0x38$	BKPT (AArch32)
$SEL_{EVT} + 0x19$	Trapped SVE	$SEL_{EVT} + 0x39$	reserved
$SEL_{EVT} + 0x1a$	Trapped ERET	$SEL_{EVT} + 0x3a$	Vector Catch (AArch32)
$SEL_{EVT} + 0x1b$	reserved	$SEL_{EVT} + 0x3b$	reserved
$SEL_{EVT} + 0x1c$	PAAuth Instruction Failure	$SEL_{EVT} + 0x3c$	BRK (AArch64)
$SEL_{EVT} + 0x1d$	reserved	$SEL_{EVT} + 0x3d$	reserved
$SEL_{EVT} + 0x1e$	reserved	$SEL_{EVT} + 0x3e$	reserved
$SEL_{EVT} + 0x1f$	reserved	$SEL_{EVT} + 0x3f$	reserved

Please refer to [3] for more details on each of these events.

*These events may be handled by the microhypervisor, in which case they will not cause portal traversals.

6.5.2 Microhypervisor Events

SEL_{OBJ}	Event
$SEL_{EVT} + SEL_{ARCH} + 0x0$	Startup
$SEL_{EVT} + SEL_{ARCH} + 0x1$	Recall
$SEL_{EVT} + SEL_{ARCH} + 0x2$	Virtual Timer


The value of SEL_{ARCH} depends on the origin of the event:

- $SEL_{ARCH} = SEL_{HST/ARCH}$ (0x40) for events that occurred in the host.
- $SEL_{ARCH} = SEL_{GST/ARCH}$ (0x40) for events that occurred in the guest.

Preliminary

6.6 Architecture-Dependent Structures

6.6.1 Hypervisor Information Page



SMG_{NUM}

Number of SMMU stream mapping groups.

CTX_{NUM}

Number of SMMU translation contexts.

6.6.2 User Thread Control Block

-		VMCR	ELRSR	+0x2d0	GIC
AP1R3	AP1R2	AP1R1	AP1R0	+0x2c0	
AP0R3	AP0R2	AP0R1	AP0R0	+0x2b0	
LR15		LR14		+0x2a0	
LR13		LR12		+0x290	
LR11		LR10		+0x280	
LR9		LR8		+0x270	
LR7		LR6		+0x260	
LR5		LR4		+0x250	
LR3		LR2		+0x240	TMR
LR1		LR0		+0x230	
CNTVOFF_EL2		CNTKCTL_EL1		+0x220	
CNTV_CTL_EL0		CNTV_CVAL_EL0		+0x210	EL2
-		HPFAR_EL2		+0x200	
FAR_EL2		ESR_EL2		+0x1f0	
SPSR_EL2		ELR_EL2		+0x1e0	
VMPIDR_EL2		VPIDR_EL2		+0x1d0	
HCRX_EL2		HCR_EL2		+0x1c0	EL1
-		MDSCR_EL1		+0x1b0	
SCTLR_EL1		VBAR_EL1		+0x1a0	
AMAIR_EL1		MAIR_EL1		+0x190	
TCR_EL1		TTBR1_EL1		+0x180	
TTBR0_EL1		AFSR1_EL1		+0x170	
AFSR0_EL1		FAR_EL1		+0x160	
ESR_EL1		SPSR_EL1		+0x150	
ELR_EL1		CONTEXTIDR_EL1		+0x140	
TPIDR_EL1		SP_EL1		+0x130	A32
-		IFSR	DACR	+0x120	
SPSR_und	SPSR_irq	SPSR_fiq	SPSR_abt	+0x110	
TPIDRRO_EL0		TPIDR_EL0		+0x100	EL0
SP_EL0		X30		+0x0f0	
X29		X28		+0x0e0	
X27		X26		+0x0d0	
X25		X24		+0x0c0	
X23		X22		+0x0b0	
X21		X20		+0x0a0	
X19		X18		+0x090	
X17		X16		+0x080	
X15		X14		+0x070	
X13		X12		+0x060	
X11		X10		+0x050	
X9		X8		+0x040	
X7		X6		+0x030	
X5		X4		+0x020	
X3		X2		+0x010	
X1		X0		+0x000	

48

32

16

0

48

32

16

0

6.6.3 Message Transfer Descriptor

The **Message Transfer Descriptor** (ARM, x86) (**MTD**), which controls the subset of the architectural state transferred during exceptions and intercepts, as described in Section 3.4.2, has the following layout:

GIC	TMR	-	EL2_HPFAR	EL2_ESR_FAR	EL2_ELR_SPSR	EL2_IDR	EL2_HCR	-	EL1_MDSCR	EL1_SCTLR	EL1_VBAR	EL1_MAIR	EL1_TCR	EL1_TTBR	EL1_AFSR	EL1_ESR_FAR	EL1_ELR_SPSR	EL1_IDR	EL1_SP	-	A32_DACR_IFSR	A32_SPSR	-	EL0_IDR	EL0_SP	FPR	GPR	ICI	POISON
31	30		27	26	25	24	23		20	19	18	17	16	15	14	13	12	11	10		8	7		5	4	3	2	1	0

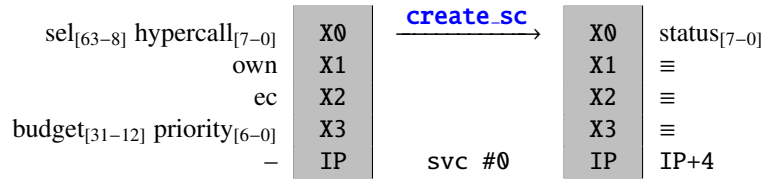
Each **MTD** bit controls the transfer of the listed architectural state to/from the respective fields in the **UTCB** (6.6.2) as follows:

- State with access r can be read from the architectural state into the **UTCB**.
- State with access w can be written from the **UTCB** into the architectural state.

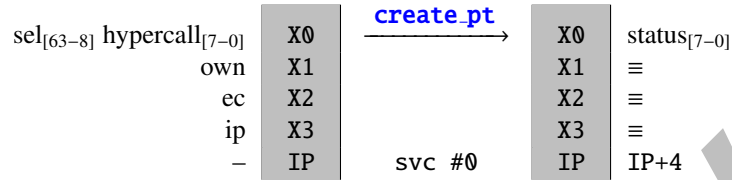
MTD Bit	Access	Host Exception State	Guest Intercept State
POISON	w	Kills the Thread	Kills the vCPU
ICI [†]	w	Invalidates the entire I-Cache	Invalidates the entire I-Cache
GPR	rw	X0 ... X30	X0 ... X30
EL0_SP	rw	SP_EL0	SP_EL0
EL0_IDR	rw	TPIDR_EL0, TPIDRRO_EL0	TPIDR_EL0, TPIDRRO_EL0
A32_SPSR	rw	-	SPSR_ABT, SPSR_FIQ, SPSR_IRQ, SPSR_UND
A32_DACR_IFSR	rw	-	DACR, IFSR
EL1_SP	rw	-	SP_EL1
EL1_IDR	rw	-	TPIDR_EL1, CONTEXTIDR_EL1
EL1_ELR_SPSR	rw	-	ELR_EL1, SPSR_EL1
EL1_ESR_FAR	rw	-	ESR_EL1, FAR_EL1
EL1_AFSR	rw	-	AFSR0_EL1, AFSR1_EL1
EL1_TTBR	rw	-	TTBR0_EL1, TTBR1_EL1
EL1_TCR	rw	-	TCR_EL1
EL1_MAIR	rw	-	MAIR_EL1, AMAIR_EL1
EL1_VBAR	rw	-	VBAR_EL1
EL1_SCTLR	rw	-	SCTLR_EL1
EL1_MDSCR	rw	-	MDSCR_EL1
EL2_HCR	rw	-	HCR_EL2, HCRX_EL2
EL2_IDR	rw	-	VPIDR_EL2, VMPIDR_EL2
EL2_ELR_SPSR	rw	ELR_EL2, SPSR_EL2	ELR_EL2, SPSR_EL2
EL2_ESR_FAR	r	ESR_EL2, FAR_EL2	ESR_EL2, FAR_EL2
EL2_HPFAR	r	-	HPFAR_EL2
TMR	rw	-	CNTV_CVAL_EL0, CNTV_CTL_EL0 CNTKCTL_EL1, CNTVOFF_EL2
GIC	rw r	-	LR0 ... LR15, APxR0 ... APxR3 ELRSR, VMCR

[†]Only affects a VIPT instruction cache of the local core. Has no effect on PIPT instruction caches, data caches, or caches of other cores.

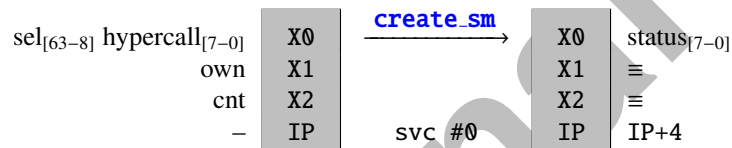
Create Scheduling Context




Create Portal



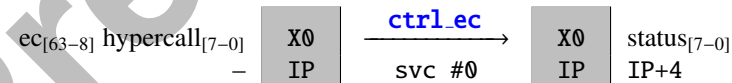
Create Semaphore



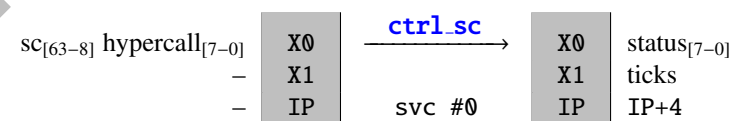
Control Protection Domain



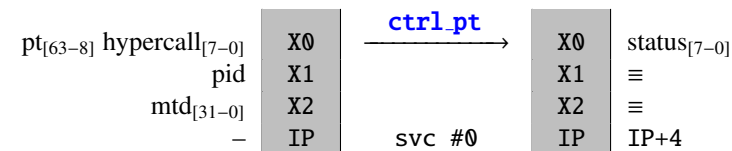
Control Execution Context




Control Scheduling Context




Control Portal




Control Semaphore




Control Power Management



Assign Interrupt



Assign Device




6.8 Supplementary Functionality

This section describes functions that do **not** conform to the calling convention for hypercalls. Because these functions cannot perform capability-based access control, their invocation is restricted to the [Root Protection Domain \(PD_{ROOT}\)](#). Invocation of these functions from any other [Protection Domain](#) generates an exception.

Secure Monitor Call

This call is proxy-filtered by the microhypervisor. If the function parameter indicates an **atomic SIP service call**, then the microhypervisor issues the corresponding SMC to the platform firmware on behalf of the caller. Otherwise this function generates an exception. Register allocation conforms to the ARM SMCCC [2].



7 ABI x86-64

7.1 Boot State

7.1.1 NOVA Microhypervisor

The bootloader must set up the [CPU](#) register state according to one of the launch types listed below when it transfers control to the NOVA microhypervisor. Furthermore, the following preconditions must be satisfied:

- The [CPU](#) state must conform to a machine state defined in the Multiboot Specification v2 [7] or v1 [6].
- All [DMA](#) activity targeting the physical memory region occupied by the microhypervisor must be quiesced. That physical memory region should also be protected against [DMA](#) accesses on systems with an IOMMU.

7.1.1.1 Multiboot v2 Launch

Only this launch type supports 64-bit [UEFI](#) platforms.

Register	Value / Description
EIP	Physical address of the NOVA Protection Domain (PD_{NOVA}) ELF image entry point
EAX	Multiboot v2 magic value (0x36d76289) [7]
EBX	Physical address of the Multiboot v2 information structure [7]
Other	~

The NOVA microhypervisor consumes the following multiboot tags, if present: 1, 3, 12, 20.

7.1.1.2 Multiboot v1 Launch

Register	Value / Description
EIP	Physical address of the NOVA Protection Domain (PD_{NOVA}) ELF image entry point
EAX	Multiboot v1 magic value (0x2badb002) [6]
EBX	Physical address of the Multiboot v1 information structure [6]
Other	~

The NOVA microhypervisor consumes the following multiboot flags, if present: 2, 3.

7.1.2 Root Protection Domain

The NOVA microhypervisor sets up the [CPU](#) register state as follows when it transfers control to the [Root Execution Context \(EC_{ROOT}\)](#):

Register	Value / Description
RIP	Virtual address of the Root Protection Domain (PD_{ROOT}) ELF image entry point
RSP	Virtual address of the Hypervisor Information Page (ARM, x86) (HIP)
RDI	EAX at boot time [†]
RSI	EBX at boot time [†]
Other	~

[†]The register contains the preserved original value from the point when control was transferred from the bootloader to the microhypervisor.

7.2 Protected Resources

The following resources are protected by the NOVA microhypervisor and are therefore [inaccessible](#) to user-mode applications.

7.2.1 Memory Space

- Physical memory occupied by the NOVA microhypervisor (conveyed via [HIP](#)).
- Physical memory occupied by Local APIC and I/O APIC devices (conveyed via [ACPI](#) MADT).
- Physical memory occupied by IOMMU devices (conveyed via [ACPI](#) DMAR or IVRS).
- Physical memory occupied by firmware runtime services (conveyed via [UEFI](#) memory map).

7.2.2 I/O Port Space

- The ACPI fixed register PM1a.CNT (conveyed via [ACPI](#) FADT).
- The ACPI fixed register PM1b.CNT (conveyed via [ACPI](#) FADT).
- The ACPI fixed register PM2.CNT (conveyed via [ACPI](#) FADT).
- The SMI_CMD port (conveyed via [ACPI](#) FADT).

7.3 Physical Memory

7.3.1 Memory Map

The [Root Protection Domain](#) ([PD_{ROOT}](#)) can obtain a list of available/reserved memory regions as follows:

- On platforms using Multiboot v2 (UEFI boot services enabled), by parsing the [UEFI memory map](#) [11].
- On platforms using Multiboot v2, by parsing the [Multiboot v2](#) memory map [7].
- On platforms using Multiboot v1, by parsing the [Multiboot v1](#) memory map [6].

7.4 Virtual Memory

The accessible virtual memory range for user-mode applications is `0 – 0x7fffffffffff`.

7.4.1 Cacheability Attributes

Encoding	ATTR_{CA}	Description
0x0	WB	Write Back
0x1	WT	Write Through
0x2	WC	Write Combining
0x3	UC	Strong Uncacheable
0x4	WP	Write Protected

Please refer to [1, 4] for details on the architectural behavior.

7.4.2 Shareability Attributes

Encoding	ATTR_{SH}	Description
0x0	UNUSED	Always use this value

7.5 Event-Specific Capability Selectors

For the delivery of exception/intercept messages, the microhypervisor performs an implicit portal traversal.

The selector for the destination portal (**SEL_{OBJ}**):

- is determined by adding the exception/intercept number to the affected EC's Event Selector Base (**SEL_{EVT}**).
- indexes into the **Object Space** (**SPC_{OBJ}**) of the affected EC's **Protection Domain** (**PD**).
- must refer to a **PT Object Capability** (**CAP_{OBJPT}**), otherwise the affected EC is killed.

7.5.1 Architectural Events

Host Exceptions

SEL_{OBJ}	Exception	SEL_{OBJ}	Exception
SEL_{EVT} + 0x0	#DE	SEL_{EVT} + 0x10	#MF
SEL_{EVT} + 0x1	#DB	SEL_{EVT} + 0x11	#AC
SEL_{EVT} + 0x2	reserved	SEL_{EVT} + 0x12	#MC*
SEL_{EVT} + 0x3	#BP	SEL_{EVT} + 0x13	#XM
SEL_{EVT} + 0x4	#OF	SEL_{EVT} + 0x14	#VE
SEL_{EVT} + 0x5	#BR	SEL_{EVT} + 0x15	#CP
SEL_{EVT} + 0x6	#UD	SEL_{EVT} + 0x16	reserved
SEL_{EVT} + 0x7	#NM*	SEL_{EVT} + 0x17	reserved
SEL_{EVT} + 0x8	#DF*	SEL_{EVT} + 0x18	reserved
SEL_{EVT} + 0x9	reserved	SEL_{EVT} + 0x19	reserved
SEL_{EVT} + 0xa	#TS*	SEL_{EVT} + 0x1a	reserved
SEL_{EVT} + 0xb	#NP	SEL_{EVT} + 0x1b	reserved
SEL_{EVT} + 0xc	#SS	SEL_{EVT} + 0x1c	reserved
SEL_{EVT} + 0xd	#GP	SEL_{EVT} + 0x1d	reserved
SEL_{EVT} + 0xe	#PF	SEL_{EVT} + 0x1e	reserved
SEL_{EVT} + 0xf	reserved	SEL_{EVT} + 0x1f	reserved

*These events may be handled by the microhypervisor, in which case they will not cause portal traversals.

†These events may be force-enabled by the microhypervisor, in which case they will cause portal traversals.

Guest Intercepts (VMX)

SEL _{OBJ}	Intercept	SEL _{OBJ}	Intercept
SEL _{EVT} + 0x0	Exception or NMI*	SEL _{EVT} + 0x28	PAUSE
SEL _{EVT} + 0x1	External Interrupt*	SEL _{EVT} + 0x29	VM Entry Failure (MCE)
SEL _{EVT} + 0x2	Triple Fault [†]	SEL _{EVT} + 0x2a	reserved
SEL _{EVT} + 0x3	INIT [†]	SEL _{EVT} + 0x2b	TPR Below Threshold
SEL _{EVT} + 0x4	SIPI [†]	SEL _{EVT} + 0x2c	APIC Access
SEL _{EVT} + 0x5	I/O SMI	SEL _{EVT} + 0x2d	Virtualized EOI
SEL _{EVT} + 0x6	Other SMI	SEL _{EVT} + 0x2e	GDTR/IDTR Access
SEL _{EVT} + 0x7	Interrupt Window	SEL _{EVT} + 0x2f	LDTR/TR Access
SEL _{EVT} + 0x8	NMI Window	SEL _{EVT} + 0x30	EPT Violation [†]
SEL _{EVT} + 0x9	Task Switch [†]	SEL _{EVT} + 0x31	EPT Misconfiguration
SEL _{EVT} + 0xa	CPUID [†]	SEL _{EVT} + 0x32	INVEPT
SEL _{EVT} + 0xb	GETSEC [†]	SEL _{EVT} + 0x33	RDTSCTP
SEL _{EVT} + 0xc	HLT [†]	SEL _{EVT} + 0x34	Preemption Timer
SEL _{EVT} + 0xd	INVD [†]	SEL _{EVT} + 0x35	INVVPID
SEL _{EVT} + 0xe	INVLPG*	SEL _{EVT} + 0x36	WBINVD, WBNOINVD
SEL _{EVT} + 0xf	RDPMSR	SEL _{EVT} + 0x37	XSETBV
SEL _{EVT} + 0x10	RDTSCTP	SEL _{EVT} + 0x38	APIC Write
SEL _{EVT} + 0x11	RSM	SEL _{EVT} + 0x39	RDRAND
SEL _{EVT} + 0x12	VMCALL	SEL _{EVT} + 0x3a	INVPCID
SEL _{EVT} + 0x13	VMCLEAR	SEL _{EVT} + 0x3b	VMFUNC
SEL _{EVT} + 0x14	VMLAUNCH	SEL _{EVT} + 0x3c	ENCLS
SEL _{EVT} + 0x15	VMPTRLD	SEL _{EVT} + 0x3d	RDSEED
SEL _{EVT} + 0x16	VMPTRST	SEL _{EVT} + 0x3e	PML Log Full
SEL _{EVT} + 0x17	VMREAD	SEL _{EVT} + 0x3f	XSAVES
SEL _{EVT} + 0x18	VMRESUME	SEL _{EVT} + 0x40	XRSTORS
SEL _{EVT} + 0x19	VMWRITE	SEL _{EVT} + 0x41	reserved
SEL _{EVT} + 0x1a	VMXOFF	SEL _{EVT} + 0x42	SPP Miss / Misconfiguration
SEL _{EVT} + 0x1b	VMXON	SEL _{EVT} + 0x43	UMWAIT
SEL _{EVT} + 0x1c	CR Access*	SEL _{EVT} + 0x44	TPAUSE
SEL _{EVT} + 0x1d	DR Access	SEL _{EVT} + 0x45	LOADIWKEY
SEL _{EVT} + 0x1e	I/O Access [†]	SEL _{EVT} + 0x46	reserved
SEL _{EVT} + 0x1f	RDMSR [†]	SEL _{EVT} + 0x47	reserved
SEL _{EVT} + 0x20	WRMSR [†]	SEL _{EVT} + 0x48	ENQCMD PASID Failure
SEL _{EVT} + 0x21	VM Entry Failure (State) [†]	SEL _{EVT} + 0x49	ENQCMD PASID Failure
SEL _{EVT} + 0x22	VM Entry Failure (MSR)	SEL _{EVT} + 0x4a	Bus Lock
SEL _{EVT} + 0x23	reserved	SEL _{EVT} + 0x4b	Notify Window
SEL _{EVT} + 0x24	MWAIT	SEL _{EVT} + 0x4c	SEAMCALL
SEL _{EVT} + 0x25	MTF	SEL _{EVT} + 0x4d	TDCALL
SEL _{EVT} + 0x26	reserved	SEL _{EVT} + 0x4e	reserved
SEL _{EVT} + 0x27	MONITOR	SEL _{EVT} + 0x4f	reserved

Please refer to [4] for more details on each of these events.

7.5.2 Microhypervisor Events

SEL _{OBJ}	Event
SEL _{EVT} + SEL _{ARCH} + 0x0	Startup
SEL _{EVT} + SEL _{ARCH} + 0x1	Recall

The value of SEL_{ARCH} depends on the origin of the event:

- SEL_{ARCH} = SEL_{HST/ARCH} (0x20) for events that occurred in the host.
- SEL_{ARCH} = SEL_{GST/ARCH} (0x100) for events that occurred in the guest.

7.6 Architecture-Dependent Structures

7.6.1 Hypervisor Information Page

The architecture-dependent [HIP](#) structure is empty.

7.6.2 User Thread Control Block

-		IA32_KERNEL_GS_BASE		+0x210
IA32_FMASK		IA32_LSTAR		+0x200
IA32_STAR		IA32_EFER		+0x1f0
IA32_PAT		IA32_SYSENTER_EIP		+0x1e0
IA32_SYSENTER_ESP		IA32_SYSENTER_CS		+0x1d0
DR7		CR8		+0x1c0
CR4		CR3		+0x1b0
CR2		CR0		+0x1a0
PDPTE3		PDPTE2		+0x190
PDPTE1		PDPTE0		+0x180
Base IDTR		Limit IDTR	-	+0x170
Base GDTR		Limit GDTR	-	+0x160
Base LDTR		Limit LDTR	AR LDTR* SEL LDTR	+0x150
Base TR		Limit TR	AR TR* SEL TR	+0x140
Base GS		Limit GS	AR GS* SEL GS	+0x130
Base FS		Limit FS	AR FS* SEL FS	+0x120
Base ES		Limit ES	AR ES* SEL ES	+0x110
Base DS		Limit DS	AR DS* SEL DS	+0x100
Base SS		Limit SS	AR SS* SEL SS	+0x0f0
Base CS		Limit CS	AR CS* SEL CS	+0x0e0
IDT Vectoring Error	IDT Vectoring Info	Interruption Error	Interruption Info [†]	+0x0d0
TPR Threshold	Exception Bitmap	PF Error Match	PF Error Mask	+0x0c0
3rd Exec Controls		2nd Exec Controls	1st Exec Controls	+0x0b0
2nd Exit Qualification		1st Exit Qualification		+0x0a0
Activity	Interruptibility	Instruction Info	Instruction Length	+0x090
RIP		RFLAGS		+0x080
R15		R14		+0x070
R13		R12		+0x060
R11		R10		+0x050
R9		R8		+0x040
R7 (RDI)		R6 (RSI)		+0x030
R5 (RBP)		R4 (RSP)		+0x020
R3 (RBX)		R2 (RDX)		+0x010
R1 (RCX)		R0 (RAX)		+0x000

*See Section 7.6.2.1 for encoding details.

[†]See Section 7.6.2.2 for encoding details.

7.6.2.1 Encoding: Segment Access Rights

~	U	G	D/B	L	AVL	P	DPL	S	Type		
	12	11	10	9	8	7	6	5	4	3	0
Field	Description										
U	0 = Segment Usable 1 = Segment Unusable										
G	Granularity										
D/B	0 = 16-bit segment 1 = 32-bit segment										
L	64-bit mode active (CS only)										
AVL	Available for use by system software										
P	Segment Present										
DPL	Descriptor Privilege Level										
S	0 = System 1 = Code or Data										
Type	Segment Type										

7.6.2.2 Encoding: Interruption Information

V	~										N	I	E	Type	Vector																					
															31											13	12	11	10	8	7					0
Field															Description																					
V															0 = Fields E, Type, Vector are invalid 1 = Fields E, Type, Vector are valid																					
N															0 = Do not request an NMI window 1 = Request an NMI window																					
I															0 = Do not request an interrupt window 1 = Request an interrupt window																					
E															0 = Do not deliver the error code from the UTCB Interruption Error field 1 = Deliver the error code from the UTCB Interruption Error field																					
Type															0 = External Interrupt 2 = Non-Maskable Interrupt 3 = Hardware Exception 4 = Software Interrupt 5 = Privileged Software Exception 6 = Software Exception 7 = Other Event (not delivered through IDT)																					
Vector															IDT Vector of Interrupt or Exception																					

7.6.3 Message Transfer Descriptor

The **Message Transfer Descriptor (ARM, x86) (MTD)**, which controls the subset of the architectural state transferred during exceptions and intercepts, as described in Section 3.4.2, has the following layout:

FPU	TLB	-	KERNEL_GS	SYSCALL	EFER	PAT	SYSENTER	DR	CR	PDPTE	IDTR	GDTR	LDTR	TR	FS/GS	DS/ES	CS/SS	INJ	STA	QUAL	CTRL	RIP	RFLAGS	GPR ₈₋₁₅	GPR ₀₋₇	POISON
31	30		23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0

Each **MTD** bit controls the transfer of the listed architectural state to/from the respective fields in the **UTCB** (7.6.2) as follows:

- State with access **r** can be read from the architectural state into the **UTCB**.
- State with access **w** can be written from the **UTCB** into the architectural state.

MTD Bit	Access	Host Exception State	Guest Intercept State
POISON	w	Kills the Thread	Kills the vCPU
GPR ₀₋₇	rw	R0 ... R7	R0 ... R7
GPR ₈₋₁₅	rw	R8 ... R15	R8 ... R15
RFLAGS	rw	RFLAGS*	RFLAGS
RIP	rw	RIP	RIP, Instruction Length, Instruction Info
CTRL	w	-	Execution Controls, Exception Bitmap, PF Error
QUAL	r	Exit Qualifications [†]	Exit Qualifications
STA	rw	-	Interruptibility State, Activity State
INJ	rw r	-	Interruption Info, Interruption Error IDT Vectoring Info, IDT Vectoring Error
CS/SS	rw	-	CS, SS (Selector, Base, Limit, AR)
DS/ES	rw	-	DS, ES (Selector, Base, Limit, AR)
FS/GS	rw	-	FS, GS (Selector, Base, Limit, AR)
TR	rw	-	TR (Selector, Base, Limit, AR)
LDTR	rw	-	LDTR (Selector, Base, Limit, AR)
GDTR	rw	-	GDTR (Base, Limit)
IDTR	rw	-	IDTR (Base, Limit)
PDPTE	rw	-	PDPTE0 ... PDPTE3
CR	rw	-	CR0, CR2, CR3, CR4, CR8
DR	rw	-	DR7
SYSENTER	rw	-	IA32_SYSENTER_{CS,ESP,EIP}
PAT	rw	-	IA32_PAT
EFER	rw	-	IA32_EFER
SYSCALL	rw	-	IA32_{STAR,LSTAR,FMASK}
KERNEL_GS	rw	-	IA32_KERNEL_GS_BASE
TLB	w	-	Invalidates the TLB for the vCPU

*Only the arithmetic flags are writable.

[†]The 1st exit qualification contains the exception error code. The 2nd exit qualification contains the fault address.

7.7 Calling Convention

The following pages describes the calling convention for each hypercall. An execution context calls into the microhypervisor by loading the hypercall identifier and other parameters into the specified processor registers and then executes the `syscall` instruction [1, 4].

The hypercall identifier consists of the hypercall number and hypercall-specific flags, as illustrated in Figure 7.1.




Figure 7.1: Hypercall Identifier


The status code returned from a hypercall has the format shown in Figure 7.2.




Figure 7.2: Status Code

The assignment of hypercall parameters to general-purpose registers is shown on the left side; the contents of the registers after the hypercall is shown on the right side.


IPC Call



IPC Reply



Create Protection Domain



Create Execution Context

sel _[63-8]	hypercall _[7-0]	RDI	→ create_ec	RDI	status _[7-0]
	own	RSI		RSI	≡
utcb _[63-12]	cpu _[11-0]	RDX		RDX	≡
	sp	RAX		RAX	≡
	evt	R8		R8	≡
	—	RCX		RCX	~
	—	R11		R11	~
	—	RIP	syscall	RIP	RIP+2

Create Scheduling Context

sel _[63-8]	hypercall _[7-0]	RDI	→ create_sc	RDI	status _[7-0]
	own	RSI		RSI	≡
	ec	RDX		RDX	≡
budget _[31-12]	priority _[6-0]	RAX		RAX	≡
	—	RCX		RCX	~
	—	R11		R11	~
	—	RIP	syscall	RIP	RIP+2

Create Portal

sel _[63-8]	hypercall _[7-0]	RDI	→ create_pt	RDI	status _[7-0]
	own	RSI		RSI	≡
	ec	RDX		RDX	≡
	ip	RAX		RAX	≡
	—	RCX		RCX	~
	—	R11		R11	~
	—	RIP	syscall	RIP	RIP+2


Create Semaphore

sel _[63-8]	hypercall _[7-0]	RDI	→ create_sm	RDI	status _[7-0]
	own	RSI		RSI	≡
	cnt	RDX		RDX	≡
	—	RCX		RCX	~
	—	R11		R11	~
	—	RIP	syscall	RIP	RIP+2


Control Protection Domain

spd _[63-8]	hypercall _[7-0]	RDI	→ ctrl_pd	RDI	status _[7-0]
	dpc	RSI		RSI	≡
src _[63-12]	ord _[6-2]	RDX		RDX	≡
dst _[63-12]	sh _[11-10]	RAX		RAX	≡
	ca _[9-7]	RCX		RCX	~
	pmm _[6-2]	R11		R11	~
	acc _[1-0]	RIP	syscall	RIP	RIP+2


Control Execution Context




Control Scheduling Context




Control Portal




Control Semaphore




Control Power Management



Assign Interrupt



Assign Device



Preliminary

Part V

Appendix

A Acronyms

ACPI	Advanced Configuration and Power Interface [12]
ATTR _{CA}	Cacheability Attribute (ARM, x86)
ATTR _{SH}	Shareability Attribute (ARM, x86)
BDF	PCI Bus:Device:Function
CAP	Capability
CAP ₀	Null Capability
CAP _{MEM}	Memory Capability
CAP _{MSR}	MSR Capability
CAP _{OBJ}	Object Capability
CAP _{OBJ_{PD}}	PD Object Capability
CAP _{OBJ_{EC}}	EC Object Capability
CAP _{OBJ_{SC}}	SC Object Capability
CAP _{OBJ_{PT}}	PT Object Capability
CAP _{OBJ_{SM}}	SM Object Capability
CAP _{PIO}	I/O Port Capability
CPU	Central Processing Unit
DMA	Direct Memory Access
EC	Execution Context
EC _{CURRENT}	Current Execution Context
EC _{ROOT}	Root Execution Context
ELF	Executable and Linkable Format [10]
FDT	Flattened Device Tree [5]
FPU	Floating Point Unit
HIP	Hypervisor Information Page (ARM, x86)
IP	Instruction Pointer
IPC	Inter-Process Communication
MSI	Message Signaled Interrupt [8]
MSR	Model-Specific Register
MTD	Message Transfer Descriptor (ARM, x86)
PCI	Peripheral Component Interconnect [8]
PD	Protection Domain
PD _{CURRENT}	Current Protection Domain
PD _{NOVA}	NOVA Protection Domain
PD _{ROOT}	Root Protection Domain
PID	Portal Identifier
PT	Portal
SC	Scheduling Context

SC_{CURRENT}	Current Scheduling Context
SC_{ROOT}	Root Scheduling Context
SEL	Capability Selector
SEL_{EVT}	Event Selector Base
SEL_{MEM}	Memory Capability Selector
SEL_{MSR}	MSR Capability Selector
SEL_{OBJ}	Object Capability Selector
SEL_{PIO}	I/O Port Capability Selector
SID	Stream Identifier
SM	Semaphore
SMMU	System Memory Management Unit
SP	Stack Pointer
SPC_{MEM}	Memory Space
SPC_{MSR}	MSR Space
SPC_{OBJ}	Object Space
SPC_{PIO}	I/O Port Space
TYPE_{SPC}	Space Type
TYPE_{ACC}	Access Type
UEFI	Unified Extensible Firmware Interface [11]
UTCB	User Thread Control Block (ARM, x86)
VMM	Virtual-Machine Monitor
ipc_call	Hypercall (ARM, x86): IPC Call
ipc_reply	Hypercall (ARM, x86): IPC Reply
create_pd	Hypercall (ARM, x86): Create Protection Domain
create_ec	Hypercall (ARM, x86): Create Execution Context
create_sc	Hypercall (ARM, x86): Create Scheduling Context
create_pt	Hypercall (ARM, x86): Create Portal
create_sm	Hypercall (ARM, x86): Create Semaphore
ctrl_pd	Hypercall (ARM, x86): Control Protection Domain
ctrl_ec	Hypercall (ARM, x86): Control Execution Context
ctrl_sc	Hypercall (ARM, x86): Control Scheduling Context
ctrl_pt	Hypercall (ARM, x86): Control Portal
ctrl_sm	Hypercall (ARM, x86): Control Semaphore
ctrl_pm	Hypercall (ARM, x86): Control Power Management
assign_int	Hypercall (ARM, x86): Assign Interrupt
assign_dev	Hypercall (ARM, x86): Assign Device


B Bibliography

- [1] *AMD64 Architecture Programmer's Manual Volume 2: System Programming*. Advanced Micro Devices, 2020. URL <https://developer.amd.com/resources/developer-guides-manuals>. Document Number: 24593. 52, 58
- [2] *ARM SMC Calling Convention*. ARM Limited, 2020. URL <https://developer.arm.com/documentation/den0028/>. Document Number: DEN0028. 49
- [3] *ARM Architecture Reference Manual ARMv8, for ARMv8-A Architecture Profile*. ARM Limited, 2021. URL <https://developer.arm.com/documentation/ddi0487/>. Document Number: DDI0487. 41, 42, 47
- [4] *Intel 64 and IA-32 Architectures Software Developer's Manual, Combined Volumes: 1, 2A, 2B, 2C, 2D, 3A, 3B, 3C, 3D and 4*. Intel Corporation, 2020. URL <https://software.intel.com/en-us/articles/intel-sdm>. Document Number: 325462. 52, 54, 58
- [5] *Devicetree Specification*. Linaro Limited, 2020. URL <https://www.devicetree.org/specifications>. Version 0.3. 39, 41, 63
- [6] Yoshinori K. Okuji, Bryan Ford, Erich Stefan Boleyn, and Kunihiro Ishiguro. *The Multiboot Specification*, 2010. URL <https://www.gnu.org/software/grub/manual/multiboot/multiboot.pdf>. Version 0.6.96. 39, 51, 52
- [7] Yoshinori K. Okuji, Bryan Ford, Erich Stefan Boleyn, Kunihiro Ishiguro, Vladimir Serbinenko, and Daniel Kiper. *The Multiboot2 Specification*, 2016. URL <https://www.gnu.org/software/grub/manual/multiboot2/multiboot.pdf>. Version 2.0. 39, 51, 52
- [8] *PCI Local Bus Specification*. PCI-SIG, 2004. URL <https://pcisig.com/specifications>. Revision 3.0. 63
- [9] Udo Steinberg and Bernhard Kauer. NOVA: A Microhypervisor-Based Secure Virtualization Architecture. In *Proceedings of the 5th ACM SIGOPS/EuroSys European Conference on Computer Systems*, pages 209–222. ACM, 2010. ISBN 978-1-60558-577-2. URL <https://doi.acm.org/10.1145/1755913.1755935>. 2
- [10] *Executable and Linking Format (ELF) Specification*. TIS Committee, 1995. URL <https://refspecs.linuxbase.org/elf/elf.pdf>. Version 1.2. 63
- [11] *Unified Extensible Firmware Interface (UEFI) Specification*. UEFI Forum, Inc., 2019. URL <https://uefi.org/specifications>. Version 2.8. 41, 52, 64
- [12] *Advanced Configuration and Power Interface (ACPI) Specification*. UEFI Forum, Inc., 2021. URL <https://uefi.org/specifications>. Version 6.4. 63

C Console

C.1 Memory-Buffer Console

The NOVA microhypervisor implements a memory-buffer console that provides run-time debug output. The memory-buffer console consists of a signaling semaphore (see 5.1.2) and an in-memory data structure with a header and a buffer as follows:



$$N = \text{MBUF End Address} - \text{MBUF Start Address} - \text{MBUF Header Size}$$

The fields of the header are used as follows:

- **RdIdx** ranges from 0 ... N-1. It points to the **next** character in the buffer that the console consumer will read and is typically advanced by the console consumer.
- **WrIdx** ranges from 0 ... N-1. It points to the **next** character in the buffer that the NOVA microhypervisor will write and is only advanced by the NOVA microhypervisor.
- The buffer is empty if **RdIdx** is equal to **WrIdx**.
- Otherwise **WrIdx** is ahead of **RdIdx**, wrapping around the buffer size N accordingly, i.e. character N+x will be stored in the same buffer slot as character x.
- If the buffer becomes full, the NOVA microhypervisor advances **RdIdx**, forcing the oldest character to be discarded from the buffer.
- At the end of each line, the NOVA microhypervisor invokes [ctrl_sm](#) (Up) on the signaling semaphore. The console consumer should use [ctrl_sm](#) (Down) on the signaling semaphore instead of polling **WrIdx**.

C.2 UART Console

Additionally several different UART consoles can be used to provide boot-time-only debug output of the microhypervisor. UART consoles should be configured for 115200 baud and 8N1 mode.

D Download

The source code of the NOVA microhypervisor and the latest version of this document can be downloaded from GitHub.

<https://github.com/udosteinberg/NOVA>

Preliminary